

Bytes, bits, etc.,

R. Inkulu

<http://www.iitg.ac.in/rinkulu/>

Outline

- 1 Self-alignment
- 2 Bitwise operators
- 3 Bit-fields
- 4 Endianness

Self-alignment of primitive types

```
int i; short s; char c; long l;

printf("%d, %d, %d, %d, %d\n", sizeof(i), sizeof(s),
      sizeof(c), sizeof(l));
//prints 4, 2, 1, 8

printf("%p, %p, %p, %p, %p\n", &i, &s, &c, &l);
//prints 0xbfd07cc, 0xbfd07ca, 0xbfd07c9, 0xbfd07d8
```


- *self-aligned*: value v of a primitive type $typeA$ is stored starting from only bytes whose address is a non-negative integer multiple of $sizeof(typeA)$
- modern processor architectures' are designed to access addresses that are self-aligned efficiently; hence, compilers generate binary code to exploit the same

Self-alignment of structures

```
typedef struct {
    char *name;
    int num;
    double price;
} Part;


int main(void) {
    Part partA;
    printf("%d, %p\n", sizeof(partA), &partA);
    //prints 16, 0xbf88bf10
}
```

- like primitive typed values, custom objects are also self-aligned: object o of a custom type (structure) T can only be stored starting from bytes whose address is a non-negative integer multiple of the most restrictive member of T ¹

¹helps in calculating internal and trailing paddings (see below) 

(Bytes, bits, etc.)

Internal padding for self-alignment



for any non-negative integer y , there must exist a non-negative integer x such that


$$4x = 8y + (50 + \delta); \text{ hence the padding } \delta \text{ before } num$$

```
typedef struct {  
    char name[50];  
    int num;  
    double price; } Part;
```

```
int main(void) {  
    Part partA;  
    printf("%d, %d, %d\n", sizeof(Part), (void*)&partA.num-(void*)&partA.name,  
        (void*)&partA.price-(void*)&partA.num); }
```

- padding is need for the purpose of member self-alignment (considering the self-alignment of the objects instantiated from that structure)

Trailing padding for self-alignment



for any non-negative integer y , there must exist a non-negative integer x such that

$8x = 8y + (62 + \delta)$; hence the padding δ at the end

```
typedef struct {  
    double price;  
    int num;  
    char name[50]; } Part;  
  
int main(void) {  
    Part partA;  
    printf("%d, %d, %d\n", sizeof(Part), (void*)&partA.num-(void*)&partA.price,  
        (void*)&partA.name-(void*)&partA.num);  
}
```

- trailing padding is needed to take care of defining array of structure objects
- however, leading padding is not allowed (according to the standard, address of first member of a structure object must need to be same as the address of object itself) (Bytes, bits, etc.,)

Examples

- `typedef struct{ short, char} A;`
- `typedef struct{ char*, int} B;`
- `typedef struct{ char*, short, int } C;`
- `typedef struct{ int, char, short} D;`
- `typedef struct{ int, char [3], short [10]} E;`
- `typedef struct { C, E [6], D [10]} F; 2`
- try more ...

homework: analyze the sizes and draw the memory layouts

homework: analyze for the best possible ordering of structure members in the structure (based on their sizes)

²recursively apply the self-alignment rules
(Bytes, bits, etc.,)

Outline

- 1 Self-alignment
- 2 Bitwise operators
- 3 Bit-fields
- 4 Endianness

Bitwise operators

- bitwise AND '&'
- bitwise OR '|'
- bitwise XOR '^'
- bitwise NOT '~'
- right shift '>>'
- left shift '<<'

Few examples

```
unsigned int i = 23;

//sets the j-th LSB of i
i = i | (1 << j-1);
//clears the j-th LSB of i
i = i & ~(1 << j-1);
//toggling the j-th LSB of i
i = i ^ (1 << j-1);

//multiplies i by 2 power j
i = i << j;
//divides i by 2 power j
i = i >> j;
//i modulo 32
unsigned int r = i & 0x1F;
//is i a power of 2
if (i & i-1 == 0) return 1;
```

Bit-masks (using macros)

in maintaining a *symbol table*, a compiler wants to categorize each identifier and/or determine the kind:

```
#define KEYWORD 01
#define EXTERNAL 02
#define STATIC 04

unsigned int flags;
...
flags |= EXTERNAL | STATIC;
    //turns on the EXTERNAL and STATIC bits in flags

flags &= ~(EXTERNAL | STATIC);
    //turns off the EXTERNAL and STATIC bits in flags

if ((flags & (EXTERNAL | STATIC)) == 0) ...
    //evaluates to true if both bits are off
```

Bit-masks (using enum)

```
enum { KEYWORD = 01, EXTERNAL = 02, STATIC = 04 };

unsigned int flags;
...
flags |= EXTERNAL | STATIC;
    //turns on the EXTERNAL and STATIC bits in flags

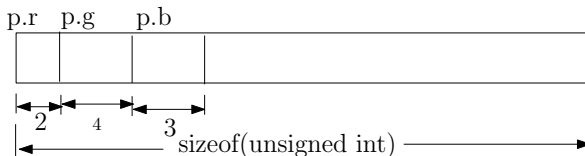
flags &= ~(EXTERNAL | STATIC);
    //turns off the EXTERNAL and STATIC bits in flags

if ((flags & (EXTERNAL | STATIC)) == 0) ...
    //evalutes to true if both bits are off
```

Outline

- 1 Self-alignment
- 2 Bitwise operators
- 3 Bit-fields**
- 4 Endianness

Intro to bit-fields



(field assignment is implementation-specific; this fig shows one such possibility)

```
struct Permissions {  
    unsigned int r : 2;  
    //only int, unsigned int types are allowed  
    unsigned int g : 4;  
    unsigned int u : 3;  
};  
struct Permissions p;    //sizeof(p) is 4
```

- contiguous set of adjacent bits is termed as a *bit-field* (in the above example, r, g and u are bit-fields)
- helps in saving space

Motivation: bit-masks using bit-fields

```
struct {  
    unsigned int isKeyword : 1;  
    unsigned int isExtern : 1;  
    unsigned int isStatic : 1  
} flags;  
  
flags.isExtern = flags.isStatic = 1;  
    //turns on both the bits  
  
flags.isExtern = flags.isStatic = 0;  
    //turns off both the bits  
  
if (flags.isExtern == 0 && flags.isStatic == 0) ...  
    //evalutes to true if both bits are off
```

- now the code is clean: bit-level optimization is not in the code instead it is hidden in the struct; together with space-efficiency

Space allocation in structures with bit-fields

```
struct Permissions {
    unsigned int r : 2;
    unsigned int g : 4;
    unsigned int u : 3;
};
int main(void) {
    struct Permissions p;
    printf("%d\n", sizeof(p));
    //prints 4
}
```

- allocates `sizeof(unsigned int)` and packs bit-fields into it successively until it is full or padding is enforced; again, allocates another `sizeof(unsigned int)` etc.,³
- bit-fields do not have addresses (hence, `&` operator not applicable)

³if the prior members are not a bit field but some other type whose size is smaller than `unsigned int`, then it allocates `sizeof(unsigned int)` to start with; when a member is not a bit-field, it utilizes the space allocated if it can accommodate it
(Bytes, bits, etc.,)

Padding for alignment

```
struct Permissions {  
    unsigned int r : 2;  
    unsigned int g : 4;  
    unsigned int u : 3;  
    unsigned int : 0;    //for padding  
    unsigned int x : 5;  
};  
  
int main(void) {  
    struct Permissions p;  
    printf("%d\n", sizeof(p));  
    //prints 8  
}
```

- unnamed fields are used for padding; special width 0 is used to force alignment at the next unsigned int boundary (next field will be stored in a new unsigned int)

Outline


- 1 Self-alignment
- 2 Bitwise operators
- 3 Bit-fields
- 4 Endianness

Description


Convention used in storing (and interpreting) bytes making an object of any primitive data type is known as *endianness*.

Two typical endian formats

little-endian representation



big-endian representation



p is the address of first byte in which an unsigned int is stored

an unsigned int: 00001010 00001011 00001100 00001101₍₂₎
(or 0x0A0B0C0D, or 168496141₍₁₀₎)

- *little-endian* machine: least significant byte of a value is stored in the smallest addressed byte of the word; ex. Intel x86 machines
- *big-endian* machine: most significant byte of a value is stored in the smallest addressed byte of the word; ex. Motorola 6800 machines

Testing endianness

```
unsigned int i = 1;
char *q = (char*)&i;

if (((int) q[0]) == 1)
    printf("little-endian\n");
else
    printf("big-endian\n");
//prints little-endian
```

Converting little-endian to big-endian (and vice versa)

```
unsigned int value = 0x0A0B0C0D, result = 0;
printf("%d bytes\n", sizeof(unsigned int));
//prints 4 bytes
char *p = (char*) &value;
printf("%d: %d, %d, %d, %d\n",
       value, p[0], p[1], p[2], p[3]);
//prints 168496141: 13, 12, 11, 10

result = (value & 0x000000FF) << 24;
result |= (value & 0x0000FF00) << 8;
result |= (value & 0x00FF0000) >> 8;
result |= (value & 0xFF000000) >> 24;
p = (char*)&result;
printf("%d: %d, %d, %d, %d\n",
       result, p[0], p[1], p[2], p[3]);
//prints 218893066: 10, 11, 12, 13
```

homework: directly modify the contents of value while using a byte of temporary variable (Bytes, bits, etc.,)

Setting a bit value in unsigned int

```
unsigned int i = 0x1;
char *p = (char*) &i;
printf("%d: %d, %d, %d, %d\n",
       i, p[0], p[1], p[2], p[3]);
//prints 1: 1, 0, 0, 0

p[1] |= 0x08;
//sets fourth least-significant bit in p[1]
printf("%d: %d, %d, %d, %d\n",
       i, p[0], p[1], p[2], p[3]);
//prints 2049: 1, 8, 0, 0
```

Binary representation of numbers

- short, int, unsigned int
- ASCII values of chars
- float⁴, double

homework:

- * given a double d , find its equivalent binary representation b
- * convert b to the relevant endian format of your machine

⁴IEEE 754 standard: <http://www.h-schmidt.net/FloatConverter/IEEE754.html>
(Bytes, bits, etc.,)

Two's complement of binary numbers

For any negative number n , the two's complement of n is the ones' complement of n added with one; for positive numbers, two's complement is same as that number itself. Advantages:

- fundamental arithmetic operations of addition, subtraction, and multiplication are identical to those for unsigned binary numbers
- zero has only a single representation
- no need to examine the signs of the operands to determine when doing addition, subtraction, and/or multiplication of numbers

based on the need, a number's two's complement is computed (with the help of hardware) just before doing the algebra in registers