

# CS245: Database Management Systems

Vijaya Saradhi

**IIT Guwahati**

Fri, 03-Apr-020

# Introduction

## Datalog

- The contents of today's self study is from the text book **Database Management Systems The Complete Book** by Hector Garcia-Molina et. al  
<https://bit.ly/2UTjzMU>
- Please refer to chapter 10 - *Logical Query Languages*
- In particular sections 10.1 and 10.2
- In this self study material and the next the idea is to introduce **Recursion** in SQL
- Simplest form of logic devised for relational model
- In its non-recursive form, datalog have the **same power** as the classical relational algebra
- By allowing recursion, we can express queries in datalog **that cannot be expressed in SQL2**
- Solution provided by datalog has been used to provide a way to allow **meaningful** recursion in the most recent SQL-99 standard

# A Logic for Relations

## Logic

- Alternative to abstract query language based on algebra, one can use a form of logic to express queries
- The logical query language **Datalog (database logic)** consists of **if-then-else** rules
- Each of these rules expresses the idea that from certain combinations of tuples in certain relations we may infer the answer to a query (that some other tuple is in some other relation)

# Predicates and Atoms

## Definitions

- Relations are represented in Datalog by predicates
- Each predicate takes a **fixed number of arguments**
- A predicate followed by its arguments is called an atom
- Syntax of atoms is just like that of function calls in conventional programming language
- Example:  $P(x_1, x_2, \dots, x_n)$  is an atom
- This atom consists of predicate  $P$  with arguments  $x_1, x_2, \dots, x_n$

# Predicates and Atoms

## Definitions

- A predicate is the name of a function
- The function that returns a **boolean value**
- If  $R$  is a relation with  $n$  attributes in some fixed order, then we shall also use  $R$  as the name of the predicate
- The atom  $R(a_1, a_2, \dots, a_n)$  have value **TRUE** if  $(a_1, a_2, \dots, a_n)$  is a tuple of  $R$
- The atom have value **FALSE** otherwise

# Predicates and Atoms

## Example

R	
A	B
1	2
3	4

- $R(1, 2)$  is true and so is  $R(3, 4)$
- For any other values  $x, y$ ,  $R(x, y)$  is false

# Predicates and Atoms

## Example

- If  $R$  is the predicate of the above example,  $R(x, y)$  is the function that tells for  $x$  and  $y$  whether the tuple  $(x, y)$  is in relation  $R$
- $R(x, y)$  returns true if  $x = 1$  and  $y = 2$ ;
- $R(x, y)$  returns true if  $x = 3$  and  $y = 4$

# Arithmetic Atoms

## Arithmetic Atoms

- $<, \leq, =, \neq, >, \geq$  are **arithmetic atoms**
- They perform comparison between two arithmetic expressions
- Both **relational atoms** and **arithmetic atoms** take as arguments values of variables
- Both return boolean values



# Datalog Rules and Queries

## Rules and Queries

- A datalog rule consists of
  - A relational atom call the **head** followed by
  - the symbol  $\leftarrow$  read as "if" followed by
  - A **body** consisting of one or more atoms known as **subgoals**

# Datalog Rules and Queries

## Example Rule

- $\text{LongMovie}(t, y) \leftarrow \text{Movie}(t, y, l, c, s, p) \text{ AND } l \geq 100$
- $t$ : title;  $y$ : year;  $l$ : length;  $c$ : inColor;  $s$ : studioName;  $p$ : producer
- The above defines **a set of long movies** which are at least 100 minutes long
- The head of the rule is the atom  $\text{LongMovie}(t, y)$
- Body consists of two subgoals

**First** Have predicate  $\text{Movie}$  with six arguments

**Arugments** a variable assuming a value;

**Interpretation**  $\text{Movie}(t, y, l, c, s, p)$  is true whenever the six variables have values that is in  $\text{Movie}$  relation

**Second**  $l \geq 100$

**Arugments** None

**Interpretation** for  $(t, y, l, c, s, p)$  which is in **Movie relation** and  $l > 100$

# Datalog Rules and Queries

## Example Rule

- When ever these variables all have values that make subgoals true, then we see what value of head is for those variables
- We add the resulting tuple to the relation whose predicate is the head

# Datalog Rules and Queries

## Restrictions

- **Safety Condition:** Every variable that appears **anywhere** in the rule **must appear in some non-negated**, relational subgoal
- This is to ensure that result of a rule is finite relation and
- Rules with arithmetic subgoals or with **negated** subgoals makes intuitive sense

# Datalog Rules and Queries

## Restrictions

- $P(x, y) \leftarrow Q(x, z) \text{ AND NOT } R(w, x, z) \text{ AND } x < y$
- Variable  $y$  appears in the head **but not in any non-negated relational subgoal**
- Variable  $w$  appears in negated, relational subgoal **but not in the non-negated, relational subgoal**
- Variable  $y$  appears in an arithmetic subgoal **but not in a non-negated, relational subgoal**

# Datalog Rules and Queries

## Example

- $P(x, y) \leftarrow Q(x, z) \text{ AND } R(z, y) \text{ AND NOT } Q(x, y)$
- Let Q consists of (1, 2) and (1, 3)
- R consists of (2, 3) and (3, 1)
- Two non-negated relational subgoals  $Q(x, z)$ ,  $R(z, y)$

$Q(x, z)$	$R(z, y)$	Consistent?	Not $Q(x, y)$ ?	Head
(1, 2)	(2, 3)	Yes	No	-
(1, 2)	(3, 1)	No; $z = 2, 3$	Irrelevant	-
(1, 3)	(2, 3)	No; $z = 3, 2$	Irrelevant	-
(1, 3)	(3, 1)	Yes	Yes	$P(1,1)$

# Extensional and Intensional Predicates

## Definitions

**Extensional Predicates (EDB)** Whose relations are **stored in a database**

**Intensional Predicates (IDB)** Whose relations are **Computed** by applying one or more Datalog rules

# Datalog Rules and Bags

## Example

$H(x, z) \leftarrow R(x, y) \text{ AND } S(y, z)$

R(A, B)		S(B,C)	
A	B	B	C
1	2	2	3
1	2	4	5
		4	5

H(x, z)	
X	Z
1	3
1	3



## Example with more than one rule

### Example

Consider the following two rules

- $H(x, y) \leftarrow S(x, y) \text{ AND } x > 1$
- $H(x, y) \leftarrow S(x, y) \text{ AND } y < 5$
- Rule 1 yields the  $H(x, y)$  set  $\{ (2, 3), (4, 5), (4, 5) \}$
- Rule 2 yields the  $H(x, y)$  set  $\{ (2, 3) \}$

# Datalog: intersection

## Definition

$I(a, b, c, d) \leftarrow R(a, b, c, d) \text{ AND } S(a, b, c, D)$

$R(a, b, c, d)$  is EDB

$S(a, b, c, d)$  is EDB

$I(a, b, c, d)$  is IDB

# Datalog: Union

## Definition

- ①  $U(a, b, c, d) \leftarrow R(a, b, c, d)$
- ②  $U(a, b, c, d) \leftarrow S(a, b, c, d)$
- ③ First rule states every tuple in R is a tuple in the IDB relation U
- ③ Second rule states every tuple in S is a tuple in the IDB relation U
- ③ Thus the two rules together represent the union

# Datalog: Difference

## Definition

$$I(a, b, c, d) \leftarrow R(a, b, c, d) \text{ AND NOT } S(a, b, c, d)$$

# Datalog: Projection

## Definition

$$I(b, d) \leftarrow R(a, b, c, d)$$

- 1 A single subgoal with predicate  $R$  is sufficient.
- 2 Arguments of this subgoal are distinct variables one for each attribute in the relation
- 3 The head has an atom with arguments that are variables corresponding to the attributes in the projection list

# Datalog: Selection

## Selection

- Simple case is when the selection condition is the AND of one more more arithmetic comparisons
- Rules are:
  - One relational subgoal for the relation upon which we are performing the selection
  - This atom has distinct variables for each component
  - one for each attribute of the relation
  - For each comparison in the selection comparison, an arithmetic subgoal that is identical to the comparison

# Datalog: Selection

## Selection Example

- $\sigma_{length \geq 100 \text{ AND } studioName = 'Fox'}(Movie)$
- $S(t, y, l, c, s, p) \leftarrow Movie(y, t, l, c, s, p) \text{ AND } l \geq 100 \text{ AND } s = 'Fox'$

# Datalog: Selection

## Selection Example

- ❶  $\sigma_{length \geq 100 \vee studioName = 'Fox'}(Movie)$
- ❷  $S(t, y, l, c, s, p) \leftarrow Movie(y, t, l, c, s, p) \text{ AND } l \geq 100$
- ❸  $S(t, y, l, c, s, p) \leftarrow Movie(y, t, l, c, s, p) \text{ AND } s = 'Fox'$ 
  - Rule 1 produces movies at least 100 minutes long
  - Rule 2 produces movies produced by 'Fox'



# Datalog: Selection

## Selection Example - Negation

- 1  $\sigma_{NOT(length \geq 100 \vee studioName = 'Fox')}(Movie)$
- 2  $\sigma_{NOT(length \geq 100) \wedge (NOT(studioName = 'Fox'))}(Movie)$
- 3  $\sigma_{(length < 100) \wedge (studioName \neq 'Fox')}(Movie)$
- 4  $S(t, y, l, c, s, p) \leftarrow Movie(y, t, l, c, s, p) \wedge l < 100 \wedge s \neq 'Fox'$

# Datalog: Cross Product

## Definition

$$I(a, b, c, d) \leftarrow R(a, b) \text{ AND } S(c, d)$$

# Datalog: Natural Join

## Definition

$$I(a, b, c, d) \leftarrow R(a, b) \text{ AND } S(b, c, d)$$

# Datalog: Theta Join

## Definition

- ①  $U \bowtie_{a < d \text{ AND } U.b \neq V.b} V$
- ②  $I(a, ub, uc, vb, vc, d) \leftarrow U(a, ub, uc) \text{ AND } V(vb, vc, d) \text{ AND } a < d$
- ③  $I(a, ub, uc, vb, vc, d) \leftarrow U(a, ub, uc) \text{ AND } V(vb, vc, d) \text{ AND } ub \neq vb$

# Datalog: Recursive Programming

## Introduction

Consider the relation

SequelOf	
movie	sequel
Rocky	Rocky II
Rocky II	Rocky III
Rocky III	Rocky IV