

CS343 - Operating Systems

Module-2C

CPU Scheduling Algorithms - 1



Dr. John Jose

Assistant Professor

Department of Computer Science & Engineering


Indian Institute of Technology Guwahati, Assam.

<http://www.iitg.ac.in/johnjose/>

Session Outline

- ❖ **CPU scheduling**
- ❖ **Categories of scheduling algorithms**
- ❖ **FCFS scheduling algorithm**
- ❖ **SJF scheduling algorithm**
- ❖ **SRTF scheduling algorithm**
- ❖ **Round Robin scheduling algorithm**
- ❖ **Priority scheduling algorithm**

Preemptive vs Non-preemptive Scheduling



1. When a process switches from the running state to the waiting state
 2. When a process switches from the running state to the ready state
 3. When a process switches from the waiting state to the ready state
 4. When a process terminates
- ❖ 1 & 4, the scheduling is **non-preemptive** (cooperative)
 - ❖ 2 & 3, the scheduling is **preemptive**

Scheduling Criteria

- ❖ Different CPU-scheduling algorithms have different properties.
- ❖ Certain characteristics/criteria are used for comparing various CPU scheduling algorithms.
 - ❖ CPU Utilization
 - ❖ Throughput
 - ❖ Turnaround time
 - ❖ Waiting Time
 - ❖ Response Time

Categories of Scheduling Algorithms & Goals

❖ Batch System

- ❖ Complete maximum number of jobs per unit time.
- ❖ Minimize time between submission and termination
- ❖ Keep CPU busy all time

❖ Interactive System

- ❖ Response to requests quickly
- ❖ Reduce waiting time
- ❖ Meet user expectations

❖ Real time System

- ❖ Meeting deadlines
- ❖ Ensure quality constraints

CPU Scheduling Algorithms

Batch Systems

- ❖ First-come first-served
- ❖ Shortest job first
- ❖ Shortest remaining Time next

Interactive Systems

- ❖ Round-robin scheduling
- ❖ Priority scheduling
- ❖ Multiple queues
- ❖ Shortest process next
- ❖ Guaranteed scheduling
- ❖ Lottery scheduling
- ❖ Fair-share scheduling

FCFS Scheduling

- ❖ Simplest CPU-scheduling algorithm
- ❖ First-come, first-served - process that requests the CPU first is allocated the CPU first
- ❖ FCFS policy is managed with a FIFO queue
- ❖ When a process enters the ready queue, its PCB is linked onto the tail of the FIFO queue
- ❖ When the CPU is free, it is allocated to process at the head of the queue
- ❖ The running process is then removed from the queue
- ❖ It is non-preemptive, once scheduled it will complete
- ❖ Short jobs wait for long

FCFS Scheduling

❖ Example: Three processes arrive in order P1, P2, P3 all at time 0.

❖ P1 burst time: 24

❖ P2 burst time: 3

❖ P3 burst time: 9



❖ Waiting Time

❖ P1: 0, P2: 24, P3: 27

❖ Completion Time:

❖ P1: 24, 2: 27, P3: 36

❖ Average Waiting Time: $(0+24+27)/3 = 17$

❖ Average Completion Time: $(24+27+36)/3 = 29$

SJF Scheduling

- ❖ Shortest (in terms of CPU time) job available is scheduled first
- ❖ Shorter processes makes progress
- ❖ SJF policy is managed with a priority queue with burst time as input
- ❖ When a process enters the ready queue, its PCB is linked onto the priority queue at the appropriate entry
- ❖ When the CPU is free, it is allocated to the process at the head of the priority queue
- ❖ If too many short jobs, long processes will starve
- ❖ SJF is non-preemptive; once allotted the process will complete
- ❖ Lowest turnaround time

SJF Scheduling

❖ Consider 3 process P2, P3, P1 all arriving at time T0.

❖ P1 burst time: 24

❖ P2 burst time: 3

❖ P3 burst time: 9



❖ Waiting Time

❖ P1: 12, P2: 0, P3: 3

❖ Completion Time:

❖ P1: 36, P2: 3, P3: 12

❖ Average Waiting Time: $(12+0+3)/3 = 5$ (compared to 17)

❖ Average Completion Time: $(36+3+12)/3 = 17$ (compared to 29)

SRTF Scheduling

- ❖ Shortest Remaining Time First (SRTF) job is scheduled first
- ❖ Preemptive scheduling algorithm
- ❖ A priority queue with remaining time is used as input
- ❖ When a process enters/re-enters the ready queue, its PCB is linked onto the priority queue at the appropriate entry
- ❖ When the CPU is free, it is allocated to the process at the head of the priority queue
- ❖ Newly arriving short process may forcefully preempt currently running process.
- ❖ Longer process may have multiple context switch before completion


SRTF Scheduling

❖ Consider the following process arriving at different time slots

❖ P1 burst time: 24 arrives at 0

❖ P2 burst time: 3 arrives at 5

❖ P3 burst time: 9 arrives at 3



❖ Waiting Time

❖ P1: $(15-3) = 12$, P2: 0, P3: $(8-5) = 3$

❖ Completion Time:

❖ P1: 36, P2: 8, P3: 15

❖ Average Waiting Time: $(0+3+12)/3 = 5$

❖ Average Completion Time: $(8+15+36)/3 = 19.6$

Round Robin Scheduling

- ❖ Modified version of preemptive FCFS
- ❖ Each process gets a small unit of CPU time (time quantum)
- ❖ FIFO queue is used as input
- ❖ When a process enters/re-enters the ready queue, its PCB is linked the tail of the queue
- ❖ After quantum expires, the process is preempted and added to the tail of the ready queue (Hence, preemptive scheduling algorithm)
- ❖ CPU is allocated to the process at the head of the queue
- ❖ Longer process may have multiple context switch before completion


Round Robin Scheduling

❖ Consider the following process arriving at T_0 , time quantum of 3 units

❖ P1 burst time: 24

❖ P2 burst time: 3

❖ P3 burst time: 9



❖ Waiting Time

❖ P1: $(6+3+3) = 12$, P2: 0, P3: $(6+3+3) = 12$

❖ Completion Time:

❖ P1: 36, P2: 3, P3: 21

❖ Average Waiting Time: $(0+12+12)/3 = 8$

❖ Average Completion Time: $(3+21+36)/3 = 20$

Round Robin Scheduling

- ❖ RR scheduling is better for short jobs and fair
- ❖ Shorter response time, good for interactive jobs
- ❖ Context-switching time adds up for long jobs
- ❖ Context switching takes additional time and overhead
- ❖ If the chosen quantum is
 - ❖ too large, response time suffers
 - ❖ infinite, performance is the same as FIFO
 - ❖ too small, throughput suffers and percentage overhead grows

Priority Scheduling

- ❖ Each process has a priority number
- ❖ Highest priority process is scheduled first; if equal priorities, then FCFS
- ❖ Managed with a priority queue with priority value as input
- ❖ When a process enters the ready queue, its PCB is linked onto the priority queue at the appropriate entry
- ❖ CPU is allocated to the process at the head of the queue
- ❖ It can have 2 variants; non-preemptive and preemptive
- ❖ Arrival of a new process with a higher priority can preempt the currently running process.

Thank you

johnjose@iitg.ac.in

<http://www.iitg.ac.in/johnjose/>

