

CS343 - Operating Systems

Module-2F

Thread Libraries



Dr. John Jose

Assistant Professor

**Department of Computer Science & Engineering
Indian Institute of Technology Guwahati, Assam.**

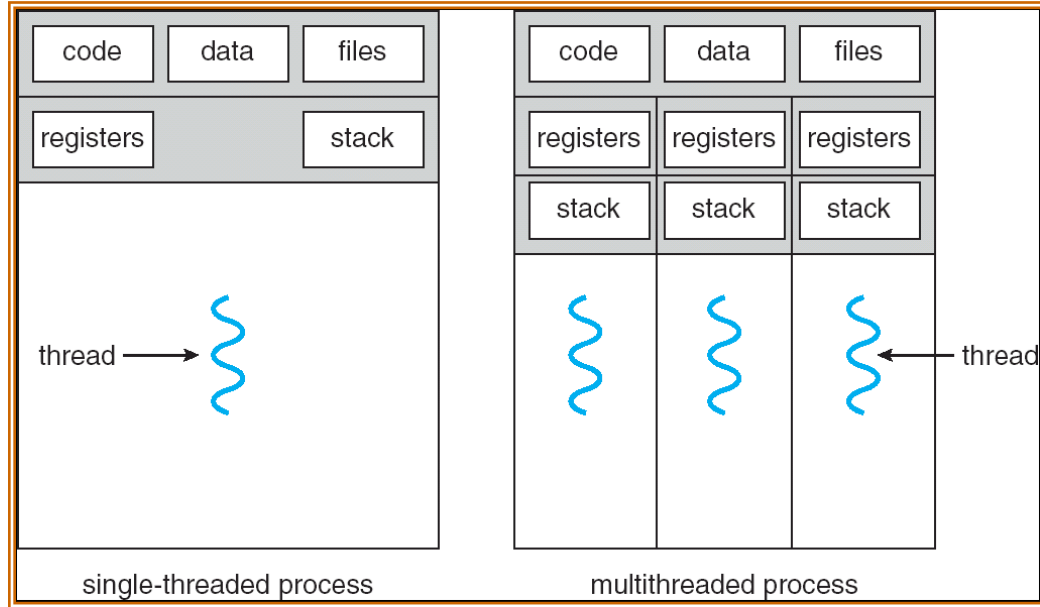
<http://www.iitg.ac.in/johnjose/>

Session Outline





- ❖ Review of Thread model
- ❖ Light weight Process
- ❖ Thread libraries
- ❖ Semantics of `fork()` and `exec()` system calls
- ❖ Thread cancellation
- ❖ Signal handling
- ❖ Thread pools
- ❖ Scheduler activations

Concept of Threads

❖ Thread is a flow of control within a process.



Thread Mappings Models



Thread Libraries

- ❖ A **thread library** provides the programmer an API for creating and managing threads
- ❖ Two primary ways to implement
 - ❖ **to provide a library entirely in user space with no kernel support.**
 - ❖ All code and data structures for the library exist in user space
 - ❖ Every function call executes in user mode, not in kernel mode
 - ❖ **to implement a kernel-level library supported by OS**
 - ❖ All code and data structures exist in kernel space
 - ❖ Invoking functions result in a system call to the kernel

Thread Libraries

- ❖ A **thread library** provides the programmer an API for creating and managing threads
- ❖ **Three** main thread libraries
 - ❖ **POSIX Pthreads** - Solaris, Linux, Mac OS, Tru64 UNIX
 - ❖ **Win32 Thread** - Windows
 - ❖ **Java Thread** - Java

Pthreads

- ❖ A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- ❖ API specifies behavior of the thread library
- ❖ Implementation is up to development of the library
- ❖ Common in UNIX operating systems (Solaris, Linux, Mac OS X)

Windows XP Threads

- ❖ Implements the one-to-one mapping
- ❖ Each thread contains a thread context that consists of
 - ❖ A thread id
 - ❖ Register set
 - ❖ Separate user and kernel stacks
 - ❖ Private data storage area

Linux Threads



- ❖ Linux refers to tasks rather than threads
- ❖ Thread creation is done through **clone()** system call
- ❖ **clone()** allows a child task to share the address space of the parent task

Java Threads


- ❖ Java threads are managed by the JVM
- ❖ Java threads may be created by:
 - ❖ Extending Thread class
 - ❖ Implementing the Runnable interface
- ❖ Because JVM is running on top of a host OS, the Java thread API is implemented using a thread library available on the host system.
 - ❖ On Windows system: using Win32 thread library
 - ❖ On Linux and Unix system: using Pthread library

Light Weight Processes (LWP)

- ❖ Most popular mapping model - **many-to-many** or **two-level** mode
- ❖ Light Weight Process (LWP)
 - ❖ An intermediate data structure between user and kernel threads
 - ❖ A user thread is attached to a LWP
 - ❖ Each LWP is attached to a kernel thread
 - ❖ OS schedules the kernel threads (not processes) to run on CPU
 - ❖ If a kernel thread blocks, LWP blocks, and the user thread blocks



Light Weight Processes (LWP)



- ❖ To the user thread library, LWP appears to be a virtual CPU on which the application can schedule a user thread to run.
- ❖ The user thread library is responsible for scheduling among user threads to the LWPs. It is similar to the CPU scheduling in kernel.
- ❖ In general, context switching between user threads involves taking a user thread of its LWP and replacing it with another thread.

Threading Issues

- ❖ Semantics of `fork()` and `exec()` system calls
- ❖ Thread cancellation
- ❖ Signal handling
- ❖ Thread pools
- ❖ Scheduler activations

Semantics of fork() and exec()

- ❖ **fork()** - system call is used to create a duplicate process.
- ❖ **exec()** - system call is used to create a new separate process.
- ❖ Fork() starts a new process which is a copy of the one that calls it, while exec replaces the current process image with another new one.
- ❖ Both parent and child processes are executed simultaneously in case of fork()
- ❖ In exec() control never returns to the original program unless there is an exec() error.

Thread Cancellation

- ❖ Terminating a thread before it has finished by other threads
 - ❖ Ex, multiple threads search DB, one thread returns result. The remaining thread might be canceled.
- ❖ Two general approaches to cancel the target thread
 - ❖ **Asynchronous cancellation** terminates the target thread immediately
 - ❖ **Deferred cancellation** allows the target thread to periodically check if it should be cancelled

Thread Cancellation

❖ Asynchronous cancellation

- ❖ The difficulty with cancellation occurs in situation where
 - ❖ resources have been allocated to a canceled thread, or
 - ❖ where a thread is canceled while in the midst of updating data it is sharing with other threads

❖ Deferred cancellation

- ❖ One thread indicates that target thread is to be canceled.
- ❖ Cancellation occurs only after the target thread has checked a flag to determine if it should be canceled or not

Signal Handling

- ❖ **Signals** are used in UNIX systems to notify a process that a particular event has occurred
- ❖ **Two types of signals**
 - ❖ Synchronous signals [illegal memory access, division by 0]
 - ❖ Asynchronous signals [Specific keystrokes (Ctrl-C), timer expire]
- ❖ What happen when a signal generated?
 - ❖ Signal is generated by particular event
 - ❖ Signal is delivered from kernel to a process
 - ❖ Signal is handled

Signal Handling

- ❖ Every signal may be handled by one of two possible handlers.
 - ❖ A default signal handler
 - ❖ A user-defined signal handler
- ❖ Every signal has a default signal handler that is run by the kernel
- ❖ The default action can be overridden by a user-defined signal handler
- ❖ Options when a signal occurs on a multi-threaded process
 - ❖ Deliver the signal to the thread to which the signal applies
 - ❖ Deliver the signal to every thread in the process
 - ❖ Deliver the signal to certain threads in the process
 - ❖ Assign a specific thread to receive all signals for the process

Thread Pools

- ❖ **Pool** of threads where they await work
- ❖ A process creates few threads at start up and place into a pool
- ❖ When receiving a request, server awakens a thread from the pool and passes the request to service
- ❖ Once the thread completes its service, it returns to the pool and await more work.
- ❖ If the pool contains no available thread, the server waits until one becomes free.
- ❖ Thread pools are faster to service a request with an existing thread than create a new thread
- ❖ Allows the number of threads in the application(s) to be bound to the size of the pool

Scheduler Activations

- ❖ Both Many-to-Many and Two-level models require communication to maintain the appropriate number of kernel threads allocated to the application
- ❖ Scheduler activations provide upcalls - a communication mechanism from the kernel to the thread library
- ❖ This communication allows an application to maintain the correct number kernel threads

Summary

- ❖ A thread is a flow of control within process.
- ❖ Multithreaded process contains several different flows of control within the same address space.
- ❖ Benefits of multi-threading includes
 - ❖ Increased responsiveness, Resource sharing within the process
 - ❖ Economy, Ability to take advantage of multiprocessor architecture
- ❖ User-level thread are thread that are visible to the programmer and are unknown to the kernel.

Summary

- ❖ OS kernel supports and manages kernel-level threads
- ❖ Three types of models relates user and kernel threads
 - ❖ One-to-one, many-to-one, many-to-many
- ❖ Thread libraries provide the application programmer with an API for creating and managing threads
 - ❖ POSIX Pthreads, Win32 threads, Java threads
- ❖ Multithreaded programs introduces several issues
 - ❖ fork()/exec(), thread cancellation, signal handling, thread pools and schedule activation.

Thank you

johnjose@iitg.ac.in

<http://www.iitg.ac.in/johnjose/>

