

CS343 - Operating Systems

Module-3D

Process Synchronization – Semaphores & Monitors



Dr. John Jose

Assistant Professor

Department of Computer Science & Engineering

Indian Institute of Technology Guwahati, Assam.

<http://www.iitg.ac.in/johnjose/>

Session Outline

- ❖ **The Critical-Section Problem**
- ❖ **Semaphores**
- ❖ **Monitors**
- ❖ **Implementation of Semaphores and Monitors**

Objectives of Process Synchronization

- ❖ To introduce the concept of process synchronization.
- ❖ To introduce the critical-section problem, whose solutions can be used to ensure the consistency of shared data
- ❖ **To present both software and hardware solutions of the critical-section problem**
- ❖ To examine several classical process-synchronization problems
- ❖ To explore several tools that are used to solve process synchronization problems

Critical Section

- ❖ Each process must ask permission to enter **critical section** in **entry section**, may follow **critical section** with **exit section**, then **remainder section**

- ❖ General structure of process P

do {

entry section

critical section

exit section

remainder section

} while (true);

```
do {  
    while (turn == j);  
        critical section  
    turn = j;  
        remainder section  
} while (true);
```

Mutual Exclusion :: Progress :: Bounded Waiting

Semaphore

- ❖ Synchronization tool for processes to synchronize their activities.
- ❖ Semaphore **S** – integer variable
- ❖ Can only be accessed via two indivisible (atomic) operations

```
wait(S)
```

```
{ while (S <= 0)  
    ; // busy wait  
    S--;  
}
```

```
signal(S)
```

```
{  
    S++;  
}
```

Semaphore Usage

- ❖ **Binary semaphore** – value can range only between 0 and 1
 - ❖ Represents single access to a resource
- ❖ **Counting semaphore** – integer value (unrestricted range)
 - ❖ Represents a resource with N concurrent access
- ❖ Consider P_1 and P_2 that require S_1 to happen before S_2
 - ❖ Create a semaphore “**synch**” initialized to 0

P1:

S_1 ;

signal(synch);

P2:

wait(synch);

S_2 ;

Semaphore Implementation

- ❖ With each semaphore there is an associated waiting queue
- ❖ Two operations:
 - ❖ **block** – place the process invoking the operation on the appropriate waiting queue
 - ❖ **wakeup** – remove one of processes in the waiting queue and place it in the ready queue

Semaphore Implementation

- ❖ Semaphore uses two atomic operations
- ❖ Each semaphore has a queue of waiting processes
- ❖ When `wait()` is called by a thread:
 - ❖ If semaphore is open, thread continues
 - ❖ If semaphore is closed, thread blocks on queue
- ❖ When `signal()` opens the semaphore:
 - ❖ If a thread is waiting on the queue, the thread is unblocked
 - ❖ If no threads are waiting on the queue, the signal is remembered for the next thread

```
wait(S)
```

```
{ while (S <= 0)  
    ;// busy wait  
    S--;  
}
```

```
signal(S)
```

```
{  
  
    S++;  
}
```


Semaphore Implementation

wait(semaphore *S)


```
{  S->value--;  
    if (S->value < 0)  
    {  
        add this process to  
        S->list;  
        block();  
    }  
}
```

signal(semaphore *S)

```
{  S->value++;  
    if (S->value <= 0)  
    {  
        remove a process P  
        from S->list;  
        wakeup(P);  
    }  
}
```

Semaphore Implementation

```
struct Semaphore {  
    int value;  
    Queue q;  
} S;  
  
withdraw (account, amount) {  
    wait(S);  
    balance = get_balance(account);  
    balance = balance - amount;  
    put_balance(account, balance);  
    signal(S);  
    return balance;  
}
```




Monitors

- ❖ A monitor is a programming language construct that controls access to shared data
- ❖ Synchronization code added by compiler, enforced at runtime
- ❖ A monitor is a module that encapsulates
 - ❖ Shared data structures
 - ❖ Procedures that operate on the shared data structures
 - ❖ Synchronization between concurrent procedure invocations
- ❖ A monitor protects its data from unstructured access
- ❖ It guarantees that threads accessing its data through its procedures interact only in legitimate ways

Monitors

- ❖ A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- ❖ Abstract data type, internal variables only accessible by code within the procedure
- ❖ One process may be active within the monitor at a time

```
monitor monitor-name
{ // shared variable declarations
  procedure P1 (...) { ... }
  procedure Pn (...) {.....}
  Initialization code (...) { ... }
}
}
```




Condition Variables

❖ Two operations are allowed on a condition variable:

❖ **x.wait()** – a process that invokes the operation is suspended until **x.signal()**

❖ **x.signal()** – resumes one of processes (if any) that invoked **x.wait()**

❖ If no **x.wait()** on the variable, then it has no effect on the variable



Condition Variables Choices

- ❖ If process P invokes **x.signal()**, and process Q is suspended in **x.wait()**, what should happen next?
 - ❖ Both Q and P cannot execute in parallel. If Q is resumed, then P must wait
- ❖ Options include
 - ❖ **Signal and wait** – P waits until Q either leaves the monitor or it waits for another condition
 - ❖ **Signal and continue** – Q waits until P either leaves the monitor or it waits for another condition

Implementation using Monitors

```
Monitor account {  
    double balance;  
  
    double withdraw(amount) {  
        balance = balance - amount;  
        return balance;  
    }  
}
```

Threads
block
waiting
to get
into
monitor

withdraw(amount)
balance = balance - amount;

withdraw(amount)

withdraw(amount)

return balance (and exit)

balance = balance - amount
return balance;

balance = balance - amount;
return balance;

When first thread exits,
another can enter.

Thank you

johnjose@iitg.ac.in

<http://www.iitg.ac.in/johnjose/>

