

Tutorial on Supervised Learning


**Part 2 : Classification using k-Nearest Neighbor (k-NN)
(implemented in Python from scratch)**

Quick Recap


k-NN - A lazy learner !

It doesn't learn a discriminative function from the training data, but “memorizes” the training dataset instead.

Example - With given data below, classify a new point (denoted by the blue star), which can belong in either red or green class.



Assume $k = 3$. Then based on the class votes received from the 3-nearest neighbors, Blue Star will be classified into the majority class.



Pre-requisite

- Euclidean distance between two points
- How **voting** works
- Dataset knowledge : **Titanic dataset** (used previously in the tutorial on Naive Bayes)

- With given dataset: Memorize a person's *Age* and *Fare*; supervision (label): he/she *survived* or not.
- Classifying a new datapoint: Using *Age* and *Fare* our learner will predict whether the person survives or not.

```
In [1]: # necessary dependencies  
  
import numpy as np  
import pandas as pd  
from collections import Counter
```

```
In [2]: # load training data

data=pd.read_csv("./train.csv")[["Survived","Age","Fare"]]
data=data.fillna(data.mean())
data.head()
```

Out[2]:

	Survived	Age	Fare
0	0	22.0	7.2500
1	1	38.0	71.2833
2	1	26.0	7.9250
3	1	35.0	53.1000
4	0	35.0	8.0500

In [3]: *# load test data*


```
passenger_id=pd.read_csv("./test.csv")["PassengerId"]  
test_data=pd.read_csv("./test.csv")[["Age","Fare"]]  
test_data=test_data.fillna(test_data.mean())  
print(test_data.head())  
test_data=test_data.values
```

	Age	Fare
0	34.5	7.8292
1	47.0	7.0000
2	62.0	9.6875
3	27.0	8.6625
4	22.0	12.2875

We need to visualize the training data to understand how k-NN works.

```
In [5]: import matplotlib.pyplot as plt
import matplotlib.patches as mpatches

col=data["Survived"]
colormap = np.array(['r', 'g'])
plt.scatter( x=data["Fare"],y=data["Age"], c=colormap[col])
surv = mpatches.Patch(color='g', label='Survived')
notSurv = mpatches.Patch(color='r', label='Did not survive')
plt.xlabel('Fare')
plt.ylabel('Age')
plt.legend(handles=[surv,notSurv])
plt.show()
```



```
In [6]: def compute_distances_one_loop(new_points, X_train):  
    #X_Train = all of our given training samples,  
    # new_points = our new points on which we want to predict  
  
    num_test = len(new_points)  
    num_train = X_train.shape[0]  
    dists = np.zeros((num_test, num_train))  
    for i in range(num_test):  
  
        # for each test point, computing Euclidean distance from all points in X  
        _train  
        difference = new_points[i] - X_train  
        difference = np.square(difference)  
        sum1 = np.sum(difference, axis=1)  
        dists[i] = np.sqrt(sum1)  
    return dists
```

```
In [7]: # Now that we have distance from each point, we will simply find out the distance which is the least  
def predict(dists, training_labels, k=3):  
    closest_y = []  
    rank = list(np.argsort(dists))  
    for x in range(0, k):  
        closest_y.append(training_labels[rank[x]])  
    closest_y = np.asarray(closest_y)  
    c=Counter(closest_y)  
    return (c.most_common()[0][0])
```

```
In [8]: dists=compute_distances_one_loop(test_data,data[["Fare","Age"]])
        print('Training instances: {}'.format(data.shape[0]))
        print('Testing instances: {}'.format(test_data.shape[0]))

        Results=[]
        for x in dists:
            Results.append(predict(x, data["Survived"]))
```

```
Training instances: 891
Testing instances: 418
```

```
In [9]: # storing the results
        f=open("result.csv","w")
        f.write("PassengerId,Survived")
        for i in range(0, len(Results)):
            f.write("\n")
            f.write(str(passenger_id[i])+", "+ str(Results[i]))

        f.close()
```

Let's compare the performance of our implementation with the kNN classifier of **scikit-learn**

```
In [10]: from sklearn.neighbors import KNeighborsClassifier

neigh = KNeighborsClassifier(n_neighbors=3)
neigh.fit(data[["Fare","Age"]], data["Survived"])
scikit_result=neigh.predict(test_data)

count=0
notr=[]
for i in range(0, len(scikit_result)):
    if scikit_result[i]==Results[i]:
        count+=1
    else:
        notr.append(i)

print ("TOTAL INSTANCES: {}".format(len(scikit_result)))
print ("RESULTS MATCHED B/W IMPLEMENTED KNN AND SCIKIT-LEARN INBUILT KNN: {}/{}".format(count, len(scikit_result)))
```

TOTAL INSTANCES: 418

RESULTS MATCHED B/W IMPLEMENTED KNN AND SCIKIT-LEARN INBUILT KNN: 416/418

End of Part 2