

CS528

High Performance Computing

A Sahu
Dept of CSE, IIT Guwahati

Outline

- Quick looks
 - HPC Trends and study requirement
 - HPC Course : Books and Contents
- Multiprocessor Intro
- Serial Code Optimization

Param Ishan HPC

- HPC system : Data Center
- Many Racks, Many rack server in a rack
- Nodes are Multicore : One rack server



Node/RackServer

Param Ishan SC @IITG

- Login Node:
 - 2x CPU login nodes, 1x GPU login node , 1x MIC login node
- Head and Management:
 - 1 pair head node (in redundant mode), 1x Management Node
- Compute Node:
 - 126x compute nodes
 - 4x high memory compute nodes
 - 16x CPU-GPU hybrid compute nodes
 - 16x CPU-MIC hybrid compute nodes.
- Network: FDR InfiniBand network
- Storage :
 - 150TB high throughput scratch space.
 - 100TB high throughput home area
 - 50TB archival for long term data storage.

Trend of HPC

- **HPC system**
 - Multi Nodes/Computer/Blades
 - **Programming Model MPI**
- **Nodes are Multicore**
 - Node have accelerators
 - **Programming Model : OpenMP, OpenCL/Cuda**
- **Core**
 - Multi Threaded, With vector instructions
 - 4 issue OOO Pipelines, Multilevel Caches,
 - **Programming Model: gcc optimized, vectorized code, OpenMP**

Need to study in HPC: User Prospects

- Single Processor
 - Architecture: Core Pipeline, Core Multithreading, Cache Hierarchy, SIMD
 - C/C++ Optimization Methods: gcc, OpenMP, Simidization, cache optimized code
- Multicore node
 - Multicore, Accelerator, Interconnections
 - OpenMP Model, Cuda Model, Accelerated Model
- HPC Server
 - Multiple Nodes/Blades, Interconnection, Storages
 - Programming Model : MPI

HPC course




- 1st Half : Before Mid Sem
 - HPC, Architecture, Programming, Code Optimization
 - Scheduling, Energy Efficiency, Power Efficiency
- 2nd Half : After Mid Sem
 - Cloud Computing
 - Mobile Cloud, FOG, EGDE, IoT

Books : Text


- Hager G and Wellein G . ***Introduction to High Performance Computing for Scientists and Engineers*** (1st ed.). CRC Press,, India, 2010.
- **Some user manuals**
- **Some recent papers**

Multiprocessor
(Will revisit again)

Example of Multicore



Tilera64



Mobile SoC Example

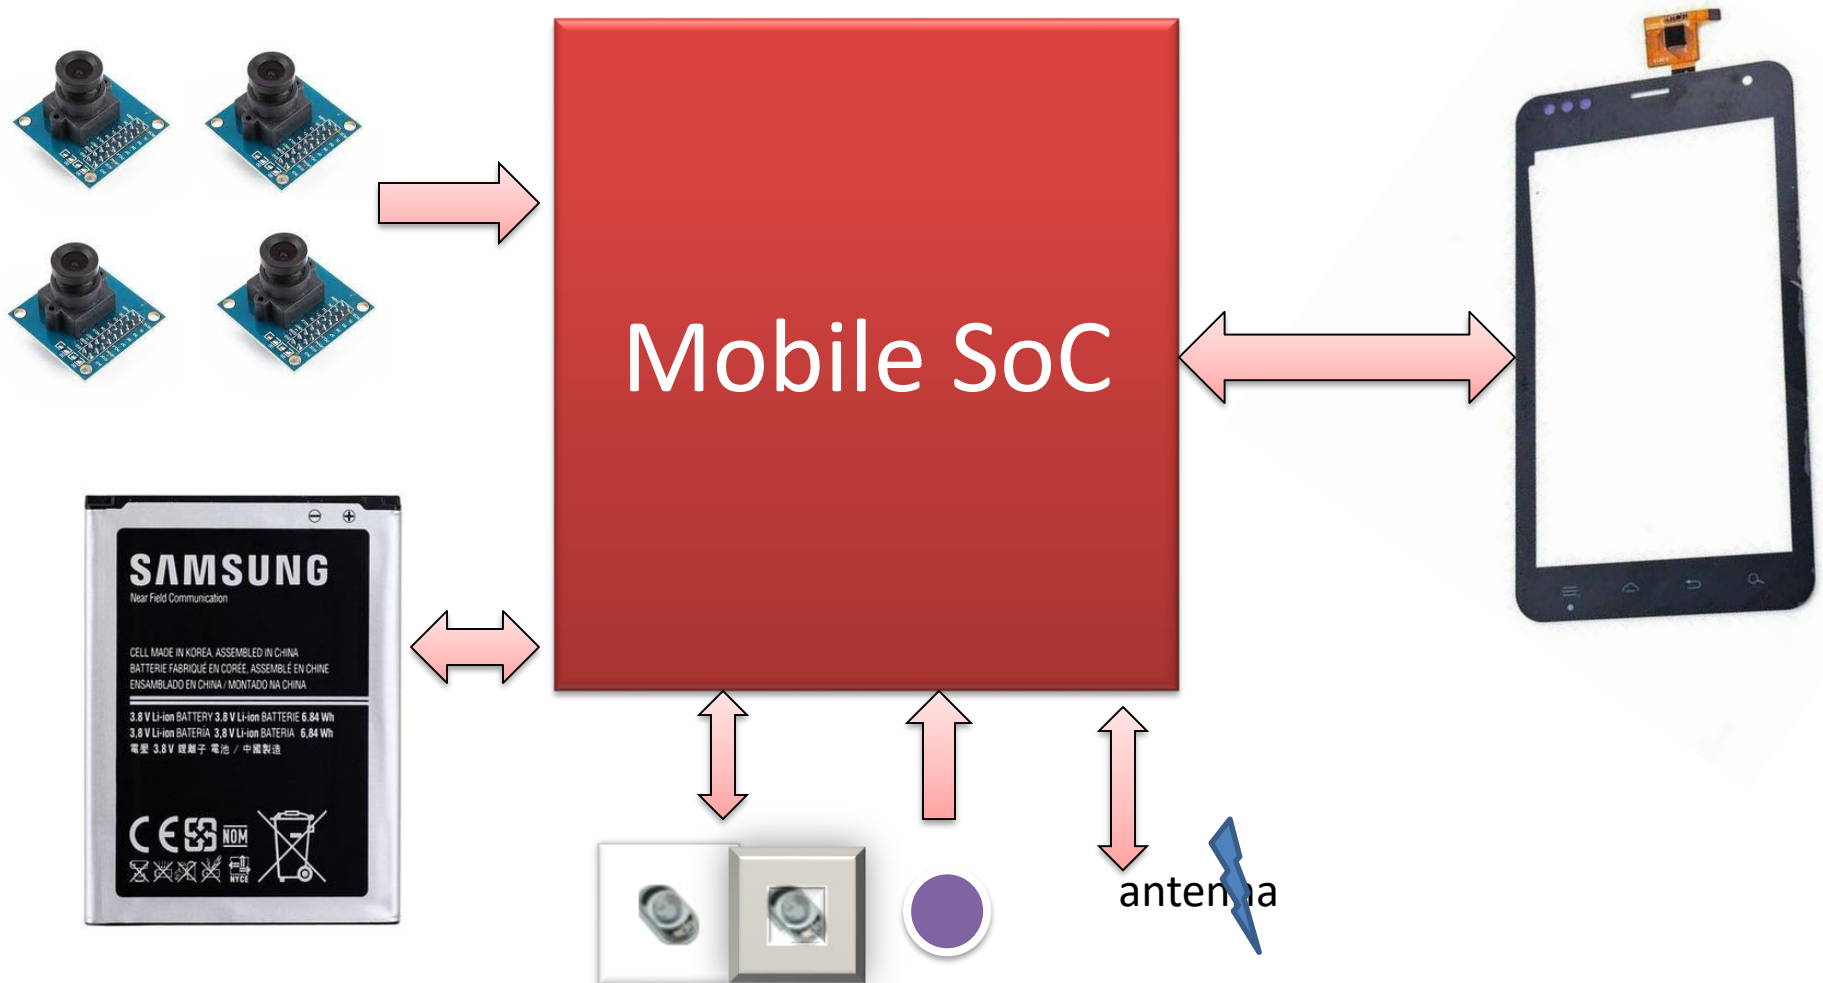
- Heterogeneous
- Diff H/W for different purpose
- Efficiency in terms
 - Perf.
 - Energy
- All in one Chip



Mobile SoC + Peripherals

Similar to motherboard and components assembly

For every components we get dozens of variety to choose



Example of Multicore Arm SoC

- Apple: A15, M1, M1X, M2
 - 2x 3.23GHz (Firestorm) + 4x 2GHz (Icestorm) or 8 core, Neural engine, GPU
- **Qualcomm: SD888, SD870**
 - y. It has 1 KryoX1@2.8, 3 A78@2.4, 4 [A55@1.8](#), AI, 5G, GPU
- Samsung: Exynos 9611
 - 4 [A73@2.3Ghz](#), 4 [A53@1.7Ghz](#), Mali G72, 5G, Codecs
- Huawei : hisilicon9000
 - 1 [A77@3.13](#), 3 [A77@2.54](#), 4 [A55@2.0](#), Mali MP24, AI, 5G, neural
- **Mediatek : Dimensity 1200**
 - 1 [A78@3.0](#), 3 [A78@2.6](#), 4 [A55@2.0](#), Mali MP24, 5G, AI,
- **Benchmarking: Antutu9, Geekbench 5, 3D Mark, Gaming FPS**

Why Multicore ?

- **Saturation of single processor performance**
 - Speed limit not to crosses : 4GHz
 - The ultimate point
 - Power consumption is proportional to cube of frequency
- $$P = \frac{1}{2} \cdot C \cdot V^2 \cdot f = k \cdot f^3$$
- as V is proportional to f*
- **Single-processor**
 - Branch prediction accuracy gone upto 95%
 - L1 Cache hits gone upto 80%
 - **ILP exploited by uniprocessor is upto 8 (mostly)**
 - **Thread/Data level parallelism needs to exploit**

Power Aware Scheduling


- $P = 1/2 C V^2 F = kF^3$ *As V-F Pairs, $V \propto F$*
 - Running Processor at 3 Ghz will consumes 27 times higher Power as compared to running at 1Ghz
- $E = kF^3 * T$
- Running a task at F and F/3 :Assume 3Ghz and 1Ghz
 - $E_f = k F^3 T$
 - $E_{f/3} = k (F/3)^3 T * 3 = k F^3 T / 9 = E_f / 9$
 - 2 times slower but 4 times energy efficient
 - 3 times slower but 9 times energy efficient
- If time permit reduce the frequency
- If task have enough slack before deadline reduce frequency

Why Multicore ?

- Application specific IC (ASIC)
 - High performance, low power than **Processor**
 - **But complexity of ASIC design is very High**
 - Example: 50MP+UHDVideo, GPS Camera in side mobile handset
 - **It is fixed for an application**

Why Multicore ?

- VLSI technology offering high integration density
- Moore's Law (In 1965, Gordon Moore Prediction)
 - Exponential growth of the number of transistors on an IC
 - Doubled every 26 months for the past three decades
- Why more transistors per IC? Smaller transistors, Larger dice



Why Multicore ?

- Many applications are highly parallel
 - Take benefit of all parallelism (instruction, data and thread)
- Multiprocessors
 - Flexible, programmable, high performance
 - Take benefit of all parallelism (instruction, data and thread)
 - Likely to be cost/power effective solutions



Why Multicore ?

- Multiprocessors are
 - Flexible, programmable, high performance
 - Processor are programmable as compared to ASIC
 - Flexible in terms of portability as compared to ASIC
 - Higher Performance than single processor

Why Multicore ?

- Multiprocessors are likely to be cost/power effective solutions
 - Share lots of resources
 - *Personal room is costlier than dormitory*
 - *You cannot allocate a Bungalow to each student: it will too costly*
 - *Hostel room with shared facility is sufficient*
 - Need not require very high frequency to run
 - Lots of replication makes easy to manage and cost effective in design

Multicore Difficulties I

- Multiprocessors are likely to be cost/power effective solutions
 - Because it share lots of resources
 - ***Personal room is costlier than dormitory***
 - Sharing resource arise many other problems
 - Critical Sections
 - Lock and Barrier Design
 - Coherence
 - Shared data at all placed should be same
 - Consistency
 - Order should be similar to serial (ROB)
 - One processor Interference others
 - Share efficiently using some policy

Multicore Difficulties II

- Many applications are highly parallel
 - Take benefit of all parallelism (instruction, data and thread)
 - Most of the coder write sequential code
 - Who will extract parallelism from applications ?
 - There is no successful auto-parallelisation tool till date
 - » Attempts: Cetus, SUIF, SolarisCC
- ***Good news: CNN/DNN python parallel library is quite successful in GPU domain***

Multicore Difficulties III

- Task scheduling in multiprocessors
 - Deterministic task scheduling on multiprocessor with more than 2 processor is NP-Complete problem
 - 4 Tasks (A,B, C and D), 3 Processors
 - {A,B,C,D}{}{} , {A,B,C}{D}{} ,Exponential Number of Solutions
- Simple Example
 - 8 tasks with execution 2, 4, 8, 5, 6, 4, 3, 20
 - Need to executed non-pre-emptively on two processor P1 and P2
 - So that overall execution time is minimized
 - **Solution : Divide 8 tasks in to two subsets, with difference of their sum is minimized ;Subset Sum Problem**

Fundamental of Serial Code Optimization

Code Optimization & Performance

- Machine-independent opt
 - Code motion, Reduction in strength, Common sub-expression Elimination
- Machine-dependent opt
 - Pointer code, Loop unrolling, Enabling instruction-level parallelism
- Tuning: Identifying perf bottlenecks
- Understanding processor Opt
 - Translation of INS into Ops, OOO
- Branches
- Caches and Blocking
- Advice

Speed and optimization

- Programmer
 - Choice of algorithm, Intelligent coding
- Compiler
 - Choice of instructions, Moving code, Reordering code, Strength reduction, *Must be faithful to original program*
- Processor
 - Pipelining, Multiple FU, Memory accesses, Branches, Caches
- Rest of system : Uncontrollable, OS, Load