

CS528
High Performance Computing

SCO
and
Data Access Optimization

A Sahu
Dept of CSE, IIT Guwahati

Outline

- SCO:
 - Simple measure Large impact : simd
 - Role of Compiler
 - Aliasing, Accuracy
 - C++ Optimization
- Data Access Optimization
 - Roofline Model
 - Caching optimization
 - App classification based DA: N/N , N^2/N^2 , N^3/N^2
- ***[Ref: Hager Book, PDF uploaded to MS Team]***

Simple measures, large impact

- Avoid Branches:
 - Code Can be SIMDized by compiler/gcc
 - Effective use of pipeline for loop code
- Use of SIMD Instruction sets
 - 512 bit AVX SIMD in modern processor
 - ML/AI app use 8 bit Ops, can be speed up $512/8=64$ time by simply SIMD-AVX

Avoid Branches: can be simdized

```
for (i=0; i<N; i++)  
    for (j=0; j<N; j++) {  
        if (i<j) S=1; else S=-1;  
        C[i] =C[i]+S*A[i][j]*B[i];  
    }
```



```
for (i=0; i<N; i++) {  
    for (j=0; j<i; j++)  
        C[i] =C[i] -A[i][j]*B[i];  
    for (j=i; j<N; j++)  
        C[i] =C[i] +A[i][j]*B[i];  
}
```

Use of SIMD: independent loop iteration

```
for (i=0; i<ALargeN; i++) {  
    A[i]=A[i]+B[i]*D[i];  
}
```

All iterations in this loop are independent : gcc SIMD utilize very nicely

//ML application uses 8 bit OPS, 512 bit AVX SIMD $512/8=64$ OPS can be done in parallel.

The ith iteration access : A[i], B[i], D[i]

Use of SIMD: independent loop iteration

```
for (i=0; i<N; i++) {  
    A[i]=A[i]+B[i]; //S1  
    B[i+1]=C[i]+D[i]; //S2  
}
```

The ith iteration access : A[i], **B[i], B[i+1]**, C[i] D[i]

Dependent loop iteration : i and i+1

Use of SIMD: independent loop iteration

```
A[0]=A[0]+B[0];  
for (i=0; i<N; i++) {  
    B[i+1]=C[i]+D[i]; //S2  
    A[i+1]=A[i+1]+B[i+1]; //S1  
}  
B[N]=C[N-1]+D[N-1];
```

**The ith iteration access :
A[i+1], B[i+1],C[i], D[i]**

Use of SIMD: independent loop iteration

```
for (i=0; i<N; i++) {  
    A[i]=A[i]+B[i]; //S1  
    B[i+1]=C[i]+D[i]; //S2  
}
```



```
A[0]=A[0]+B[0];  
for (i=0; i<N; i++) {  
    B[i+1]=C[i]+D[i]; //S2  
    A[i+1]=A[i+1]+B[i+1]; //S1  
}  
B[N]=C[N-1]+D[N-1];
```


Use of SIMD: independent loop iteration

- Affine access : index $a.x+b$ form

```
for (i=0; i<N; i++) {  
    X[a*i+b] = X[c*i+d];  
    //where a,b,c,d are integer  
}
```

- $\text{GCD}(c,a)$ divides $(d-b)$ for loop dependence
- Ref Book : Hennesy Paterson, Advanced Computer Architecture, 5th Edition Book
– One can refer any Adv. Compiler Book

GCD test Example

```
for (i=0; i<N; i++) {  
    X[2*i+3]=X[2*i]+5.0;  
    //X[a*i+b]=X[c*i+d]+k;  
}
```

- $\text{GCD}(c,a)$ must divide $(d-b)$ for loop dependence
- Value of $a=2$, $b=3$, $c=2$, $d=0$;
- $\text{GCD}(a,c)=2$, $d-b=-3$
- 2 does not divide -3 \rightarrow No dependence Possible

Role of Compilers

- General Compiler Optimization Options
- Inlining
- Aliasing
- Computational Accuracy

General Compiler Optimization Options

- GCC optimization : -O0, -O1, -O2,-O3
- \$man gcc
- At -O0 level:
 - Compiler refrain from most of the opt.
 - It is correct choice for analyzing the code with debugger
- At high level
 - Mixed up source lines, eliminate redundant variable, rearrange arithmetic expressions
 - Debugger has a hard time to give user a consistent view on code and data

General Compiler Optimization Options

- Level 1
 - fauto-inc-dec, -fmove-loop-invariant, -fmerge-constants, -ftree-copy-prop, -finline-fun-called-once
- Level 2
 - -falign-functions, -falign-loops, level, -finlining-small-fun, -finling-indirect-fun, -freorder-fun, -fstrict-aliasing
- Level 3
 - -ftree-slp-vectorize, -fvect-cost-model

Inlining

- Inlining
 - Tries to save overhead by inserting the complete code of function
 - At the place where it called
- Saves time and resources by
 - not using function call, stack
 - All compiler to use registers
 - Allows compiler to views a larger portion of code and employ OPTimization
- Auto inline or hint in program to function to be inlined

Aliasing

```
void scale_shift(double *a, double *b,  
                double s, int n) {  
    for(int i=1; i<n; i++)  
        a[i]=s*b[i-1];  
}
```

- Assuming a and b don't overlap
 - **double __restrict *a, double __restrict *b**
 - **__restrict say no overlap**
- Load and stores in the loop can be rearranged by compiler
- Apply software-pipelining, unrolling, group load/store, SIMD, etc

Computational Accuracy

- Compiler some time refrain from re-arranging arithmetic expression
- FP domain associative rule
 $a+(b+c) \neq (a+b)+c$
- If accuracy need to be maintained
 - Compared to non optimized code
 - Associative rules must not be used by compiler
 - Should be left to programmer to regroup safely
- FP underflow are push to zero

Computational Accuracy

- FP domain associative rule **$a+(b+c) \neq (a+b)+c$**
 - Let $a=1.0 \times 10^{38}$, $b=-1.0 \times 10^{38}$, $c=1$
 - Result of $a+(b+c)$
 - $= 1.0 \times 10^{38} + (-1.0 \times 10^{38} + 1)$
 - $= 1.0 \times 10^{38} + (-1.0 \times 10^{38}) = 0$ **//Big+Small=Big**
 - Result of $(a+b)+c$
 - $= (1.0 \times 10^{38} + -1.0 \times 10^{38}) + 1$
 - $= 0 + 1 = 1$

Computational Accuracy

- Why it happens for FP?
 - **FP format use 32 bit represent number up to $\pm 2^{127}$**
 - **Int use 32 bit represent up to $\pm 2^{31}$**
 - Used same 32 bit for large numbers, numbers are not equal-spaced
 - From 36000ft, both IITG and Amingoan are not distinguishable [Resolution:]
 - Going by Air: Delhi, Noida, Gurgaon use the same Airport

C++ Optimizations

- Temporaries
- Dynamic Memory Management
- Loop Kernel and Iterators

C++ Opt: Temporaries

- C++: operator overloading uses

```
class vec3d{
    double x,y,z;
public: vec3d( double _x=0.0, _y=0.0, _z=0.0):x(_x),y(-y),z(_z){}
    vec3d operator+(const vec3d &oth){
        vec3d tmp; tmp.x=x+oth.x; ...for y, and z
        return tmp
    }
    vec3d operator*(double s, const vec3d &v){
        vec3d tmp(s*v.x, s*v.y,s*v.z); return tmp;}
}

main() {
    vec3d a, b(2,2), c(3); double u=1.0,v=2.0;
    a=u*b + v*c;
}
```

C++ Opt: Temporaries

- C++: operator overloading uses
- In this prev statements
 - Constructor get called for a,b,c
 - Operator*, constructor for tmp, destructor for tmp
 - Operator*, constructor for tmp, destructor for tmp
 - Operator+, constructor for tmp, destructor for tmp
 - Copy constrtor called with tmp
- **Simply we could have write**
 - **`a.x=u*b.x+v*c.x; a.y=u*b.y+v*c.y; a.z=u*b.z+v*c.z;`**

C++ Opt: Dynamic Memory Management

```
void func(double Th, int Len) {  
    vector<double> v(Len);  
    if(rand() > Th * RAND_MAX) {  
        v = obtain_data(Len);  
        sort(v.begin(), v.end());  
        process_data(v);  
    }  
}
```



This creation is
Costly

C++ Opt: Dynamic Memory Management

```
void func(double Th, int Len) {  
  
    if (rand() > Th * RAND_MAX) {  
        vector<double> v(Len);  
        v = obtain_data(Len);  
        sort(v.begin(), v.end());  
        process_data(v);  
    }  
}
```



**This creation is
Costly, so make it
Lazy**

- Lazy construction : if the probability of requirement is low
 - Post pone the construction if the condition become true

C++ Opt: Dynamic Memory Management

```
void func(double Th, int Len) {  
    static vector<double> v(LargeLen);  
    if(rand() > Th * RAND_MAX) {  
        v = obtain_data(Len);  
        sort(v.begin(), v.end());  
        process_data(v);  
    }  
}
```



**One time
construction for
all calls**

- Static Construction : if the probability of requirement is high or always required
 - one time Construction : for all call/invocation
 - Take sufficient largeLen

C++ Opt: Loop Kernel and Iterators

- Runtime of scientific application dominated by loops or loops nest
- Compiler ability to optimize loops is pivotal for getting performance
- Operator overloading and template may hinders good loop optimization

C++ Opt: Loop Kernel and Iterators

- Non-SIMDized code: operator[] called twice for a and b, compiler refuse to SIMDize

```
template<class T>
T Sprod(const vector<T> &A,
        const vector<T> &B) {

    T result=T(0);
    int s=A.size();
    for(int i=0;i<s;i++)
        result += A[i]*B[i]; //Access
    return result;
}
```

C++ Opt: Loop Kernel and Iterators

- SIMDized

```
template<class T>
T Sprod(const vector<T> &A,
        const vector<T> &B) {
vector<T>::const_iterator
    iA=A.begin(), iB=B.begin();
    T result=T(0);
    int s=A.size();
    for(int i=0;i<s;i++)
        result += iA[i]*iB[i]; //Access
    return result;
}
```

Data Access Optimization

Performance of System: Modeling Customer Dispatch in a Bank

Resolving door
Throughput:
 b_s [customer/sec]



Processing
Capability:
 P_{peak} [task/sec]

Intensity:
 I [task/customer]



Modeling Customer Dispatch in a Bank

- How fast can tasks be processed? $P[\text{tasks/sec}]$
- The bottleneck is either
 - The service desks (peak. tasks/sec): P_{peak}
 - The revolving door (max. customers/sec): $I \cdot b_s$
- Performance $P = \min(P_{\text{peak}}, I \cdot b_s)$
- This is the “Roofline Model”
 - High intensity: P limited by “execution”
 - Low intensity: P limited by “bottleneck”