

CS528

**Algo/App Classification : DAO
&
Multi-threading and OpenMP**

A Sahu

Dept of CSE, IIT Guwahati

Outline

- APP classification based on Data access
- Threading
 - Pthread, Locking
- OpenMP
 - AutoParallelization
- MPI

Algorithm Classification and Access Optimization

- $O(N)/O(N)$: If the # of arithmetic Ops and data transfer (LD/ST) are proportional to Loop Length N
 - Optimization potential is limited
 - Example Scalar Product, vector add, sparse MVM
- Memory bound for large N
- Compiler generated code achieve good perf.
 - Using software pipelining and loop nests

Loop fusion for $O(N)/O(N)$

```
for (i=0; i<N; i++)  
    A[i]=B[i]+C[i];    //  $B_c=3W/1F$   
for (i=0; i<N; i++)  
    Z[i]=B[i]+E[i];    //  $B_c=3W/1F$ 
```



```
for (i=0; i<N; i++) {  
    A[i]=B[i]+C[i];  
    Z[i]=B[i]+E[i];  
}
```

$B_c=5W/2F$
No need get to
 $B[i]$ again

$O(N^2)/O(N^2)$: OPS/DataTransfer

- Typical two loop nests with loop strip count N
 - $O(N^2)$ operation for $O(N^2)$ loads and stores
- Example: dense MVM, Mat add, MatTrans
- MVM : -> Covert both access to row access

```
for (i=0; i<N; i++) {  
    tmp=C[i]  
    for (j=0; j<N; j++)    tmp=A[i][j]*B[j]  
    C[i]=tmp  
}
```

- Row I of A and vector B
- Original $Bc=2W/2F$ but $\rightarrow 2W*m/2F$
- m is miss rate of cache for Row access

$O(N^3)/O(N^2)$: OPS/DataTransfer

- Typical three loop nests
 - $O(N^3)$ operation for $O(N^2)$ loads and stores
- Example: dense Matrix Multiplication
- Implementation of cache Bound
 - Already studied : loop interchange
 - **Blocking : Strassen multiplication, will be discussed later**

Multiprocessor Programming using Threading

Threading Language and Support

- Initially threading used for
 - Multiprocessing on single core : earlier days
 - Feeling/simulation of doing multiple work simultaneously even if on one processor
 - Used TDM, time slicing, Interleaving
- Now a day threading mostly used for
 - To take benefit of multicore
 - Performance and energy efficiency
 - We can use both TDM and SDM

Threading Language and Support

- Pthread: POSIX thread
 - Popular, Initial and Basic one
- Improved Constructs for threading
 - c++ thread : available in c++11, c++14
 - **Java thread : very good memory model**
 - Atomic function, Mutex
- Thread Pooling and higher level management
 - OpenMP (loop based)
 - Cilk (dynamic DAG based)

Pthread, C++ Thread, Cilk and OpenMP

```
pthread_t tid1, tid2;
pthread_create(&tid1, NULL, Fun1, NULL);
pthread_create(&tid2, NULL, Fun2, NULL);
pthread_join(tid1, NULL);
pthread_join(tid2, NULL);
```

Pthread

```
thread t1(Fun1);
thread t1(Fun2, 0, 1, 2); // 0, 1, 2 param to Fun2
t1.join();
t2.join();
```

C++
thread

```
#pragma omp parallel for
```

```
for(i=0;i<N;i++)
```

```
    A[i]=B[i]*C[i];
```

```
//Auto convert serial code to threaded code
```

```
// $gcc -fopenmp test.c; export OMP_NUM_THREADS=10; ./a.out
```

```
cilk fib (int n){//Cilk dynamic parallism, DAG recursive code
```

```
    if (n<2) return n;
```

```
    int x=spawn fib(n-1); //spawn new thread
```

```
    tnt y=spawn fib(n-2); //spawn new thread
```

```
    sync;
```

```
    return x+y;
```

```
}
```

Programming with Threads

- Threads
- Shared variables
- The need for synchronization
- Synchronizing with semaphores
- Thread safety and reentrancy
- Races and deadlocks

Traditional View of a Process

- Process = process context + code, data, and stack

Process context

Program context:

Data registers

Condition codes

Stack pointer (SP)

Program counter (PC)


Kernel context:

VM structures (VMem)

Descriptor table

brk pointer

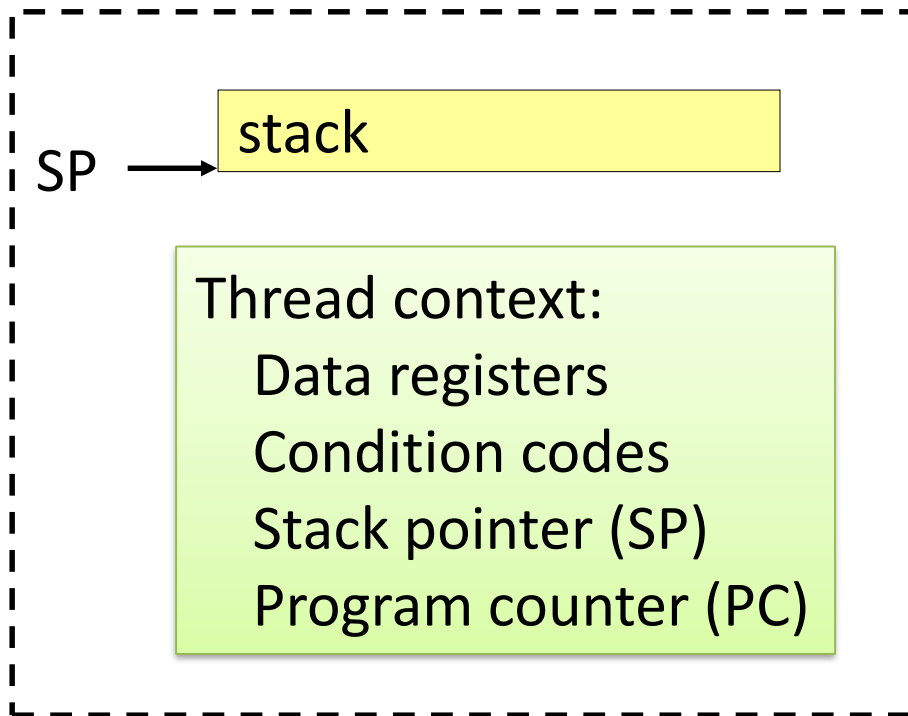
Code, data, and stack



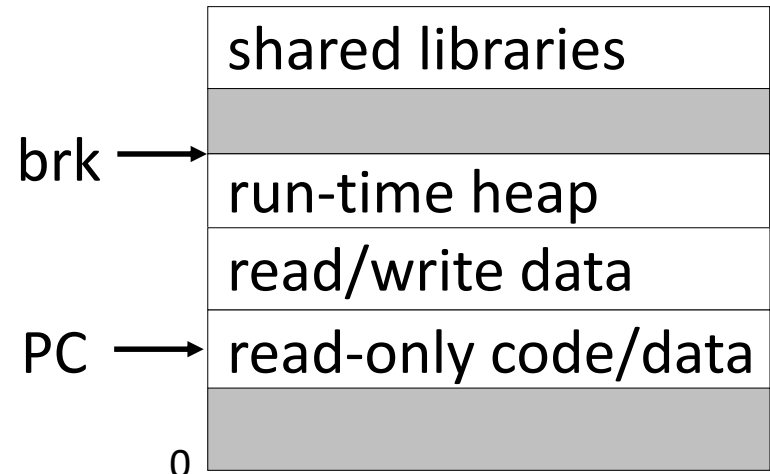
Alternate View of a Process

- Process = thread+ code, data & kernel context

Thread (main thread)



Code and Data



Kernel context:
VM structures
Descriptor table
brk pointer

A Process With Multiple Threads

- Multiple threads can be associated with a process
 - Each thread has its *own* logical control flow (sequence of PC values)
 - Each thread *shares* the same code, data, and kernel context
 - Each thread has its own thread id (TID)

A Process With Multiple Threads

Thread 1 (main thread)

stack 1

Thread 1 context:

Data registers

Condition codes

SP1

PC1

Shared code
and data

shared libraries

run-time heap

read/write data

read-only code/data

0

Kernel context:

VM structures

Descriptor table

brk pointer

Thread 2 (peer
thread)

stack 2

Thread 2 context:

Data registers

Condition codes


SP2

PC2


Logical View of Threads

- Threads associated with a process form a pool of peers
 - Unlike processes, which form a tree hierarchy

Threads associated with
process foo



Process hierarchy



Posix Threads (Pthreads) Interface

- **Creating and reaping threads**
 - `pthread_create`, `pthread_join`
- **Determining your thread ID** : `pthread_self`
- **Terminating threads**
 - `pthread_cancel`, `pthread_exit`
 - `exit` [terminates all threads], `return` [terminates current thread]
- **Synchronizing access to shared variables**
 - `pthread_mutex_init`,
`pthread_mutex_[un]lock`
 - `pthread_cond_init`,
`pthread_cond_[timed]wait`

The Pthreads "hello, world" Program

```
/* thread routine */  
void *HelloW(void *vargp) {  
    printf("Hello, world!\n");  
    return NULL;  
}
```

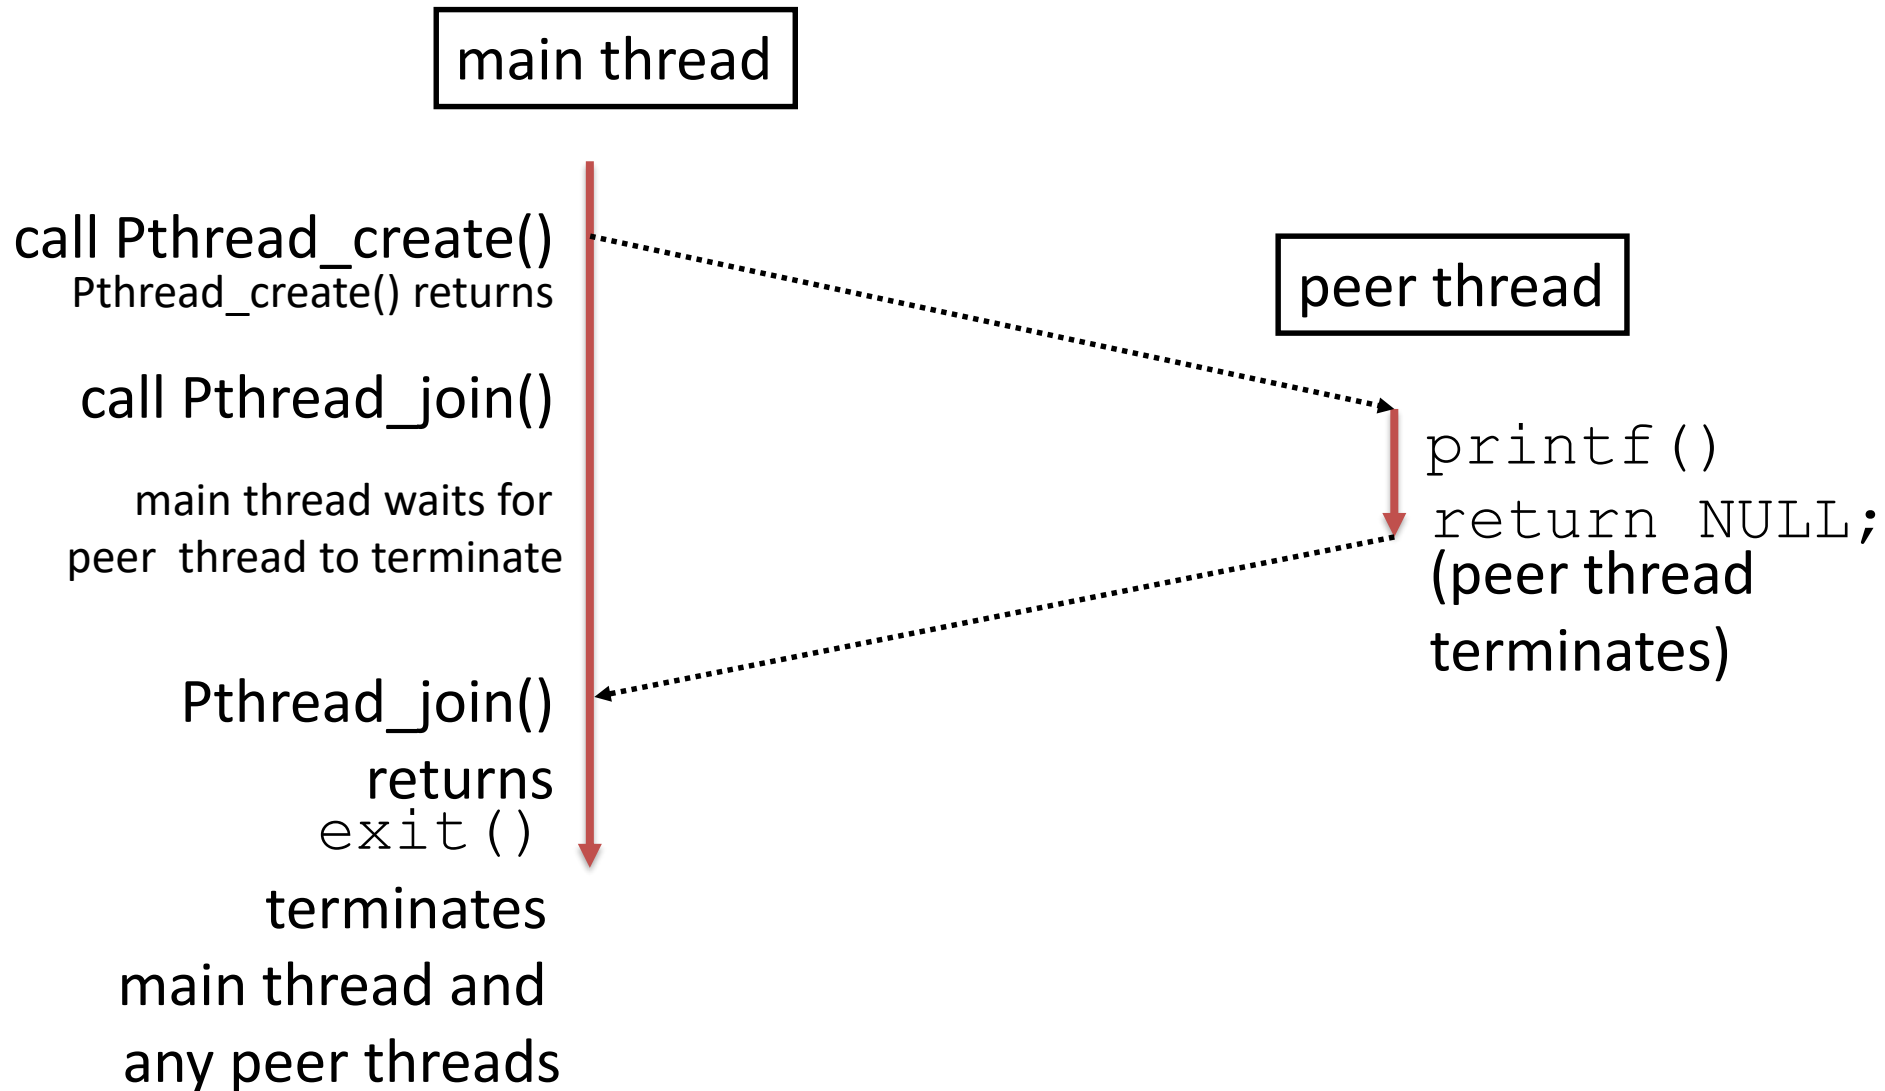
```
int main() {  
    pthread_t tid;  
    pthread_create(&tid, NULL, HelloW, NULL);  
    pthread_join(tid, NULL);  
    return 0;  
}
```

*Thread attributes
(usually NULL)*

*Thread arguments
(void *p)*

*return value
(void **p)*

Execution of Threaded “hello, world”



Pros and Cons: Thread-Based Designs

- **+ Easy to share data structures between threads**
 - E.g., logging information, file cache
- **+ Threads are more efficient than processes**
- **– Unintentional sharing can introduce subtle and hard-to-reproduce errors!**
 - Ease of data sharing is greatest strength of threads
 - Also greatest weakness!

VectorSum Serial

```
int A[VSize], B[VSize], C[VSize];

void VectorSumSerial() {
    for( int j=0; j<SIZE; j++)
        A[j]=B[j]+C[j];
}
```

Suppose Size=1000

0-249	250-499	500-749	750-999
--------------	----------------	----------------	----------------

T1

T2

T3

T4

VectorSum Serial

```
int A[VSize], B[VSize], C[VSize];

void VectorSumSerial() {
    for( int j=0; j<SIZE; j++)
        A[j]=B[j]+C[j];
}
```

- Independent
- Divide work into equal for each thread
- Work per thread: $\text{Size}/\text{numThread}$

VectorSum Parallel

```
void *DoVectorSum(void *tid) {  
    int j, SzPerthrd, LB, UB, TID;  
    TID= *((int *)tid);  
    SzPerthrd=(VSize/NUM_THREADS);  
    LB= SzPerthrd*TID;UB=LB+SzPerthrd;  
  
    for (j=LB; j<UB; j++)  
        A[j]=B[j]+C[j];  
}
```



VectorSum Parallel

```
int main() {  
    int i;  
    pthread_t thread[NUM_THREADS];  
    for (i = 0; i < NUM_THREADS; i++)  
        pthread_create(&thread[i],  
            NULL, DoVectorSum, (void*)&i);  
    for (i = 0; i < NUM_THREADS; i++)  
        pthread_join(thread[i], NULL);  
  
    return 0;  
}
```


Matrix multiply and threaded matrix multiply

- Matrix multiply: $C = A \times B$

$$C[i, j] = \sum_{k=1}^N A[i, k] \times B[k, j]$$



Matrix multiply and threaded matrix multiply

- Matrix multiply: $C = A \times B$

$$C[i, j] = \sum_{k=1}^N A[i, k] \times B[k, j]$$

- Divide the whole rows to T chunks
 - Each chunk contains : N/T rows, Assume $N\%T=0$

Matrix multiply Serial

```
void MatMul () {  
    int i, j, k, S;  
    for (i=0; i<Size; i++)  
        for (j=0; j<Size; j++) {  
            S=0;  
            for (k=0; k<Size; k++)  
                S=S+A[i][k]*B[k][j];  
            C[i][j]=S;  
        }  
}
```

Matrix Pthreaded: RowWise

```
void * DoMatMulThread(void *arg) {  
    int i,j,k,S, LB,UB, TID, ThrdSz;  
    TID=*((int *)arg);ThrdSz=Size/NumThrd;  
    LB=TID*ThrdSz;UB=LB+ThrdSz;  
  
    for (i=LB; i<UB; i++)  
        for (j=0; j<Size; j++) {  
            S=0;  
            for (k=0; k<Size; k++)  
                S=S+A[i][k]*B[k][j];  
            C[i][j]=S;  
        }  
}
```

Matrix Pthreaded: RowWise

```
int main() {  
    pthread_t  thread[NumThread];  
    int t;  
    Initialize();  
    for (t=0; t<NumThread; t++)  
        pthread_create(&thread[t], NULL,  
            DoMatMulThread, &t);  
    for (t=0; t<NumThread; t++)  
        pthread_join(thread[t], NULL);  
    TestResult();  
    return 0;  
}
```

Estimating π using Monte Carlo


- The probability of a random point lying inside the unit circle:

$$\mathbf{P} \left(x^2 + y^2 < 1 \right) = \frac{A_{circle}}{A_{square}} = \frac{\pi}{4}$$

- If pick a random point N times and M of those times the point lies inside the unit circle:

$$\mathbf{P}^{\circ} \left(x^2 + y^2 < 1 \right) = \frac{M}{N}$$

- If N becomes very large, $\mathbf{P} = \mathbf{P}^{\circ}$



$$\pi = \frac{4 \cdot M}{N}$$

Value of PI: Monte-Carlo Method

```
void MontePI () {  
    int count=0,i;  
    double x,y,z;  
    for ( i=0; i<niter; i++) {  
        x = (double)rand()/RAND_MAX;  
        y = (double)rand()/RAND_MAX;  
        z = x*x+y*y;  
        if (z<=1) count++;  
    }  
    pi=(double) count/niter*4;  
}
```

PI- Multi-threaded

- 1 thread you are able to generate N points
 - Suppose M points fall under unit circle
 - $PI = 4M/N$
- With 10 thread generate 10XN points and calculate more accurately
 - Each thread calculate own value of PI (or M)
 - Average later on (or recalculate PI from collective M)

Value of PI: Pthreaded

```
int main() {  
    pthread_t  thread[NumThread]; double pi;  
    int t, at[NumThread], count, TotalIter;  
    for(t=0; t<NumThread; t++)  
        pthread_create(&thread[t], NULL,  
                        DoLocalMC_PI, &t);  
    for(t=0; t<NumThread; t++)  
        pthread_join(thread[t], NULL);  
    for(t=0; t<NumThread; t++) count+=LCount[t];  
    TotalIter=niter*NumThread;  
    pi=((double) count/TotalIter)*4;  
    return 0;  
}
```

Value of PI: Pthreaded

```
int LCount [NumThread];  
void *DoLocalMC_Pi (void *aTid) {  
    int tid, count, i; double x, y, z;  
    tid= *((int *) aTid);  
    count=0; LCount[tid]=0;  
    for ( i=0; i<niter; i++) {  
        x = (double) rand() / RAND_MAX;  
        y = (double) rand() / RAND_MAX;  
        z = x*x+y*y; if (z<=1) count++;  
    }  
    LCount[tid]=count;  
}
```