


GPU & Application Analysis for GPGPU


A. Sahu

Dept of CSE, IIT Guwahati

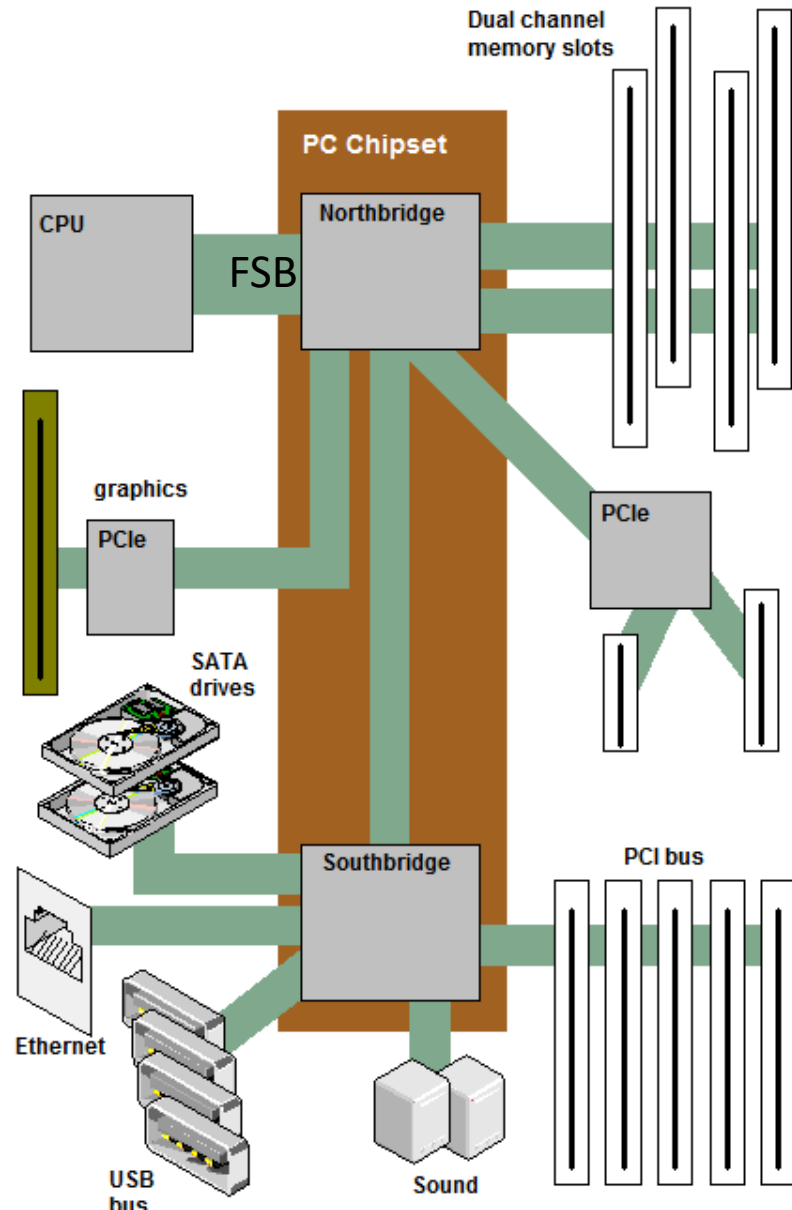

GPU

- Graphics Cards to Motherboard PCI Slot
 - Peripheral Components Interconnect
- To accelerate Graphics computation
- Earlier day : It was fixed purpose
- Now a days, it is programmable, configurable
 - Why not to use them for general purpose?
 - For what kind of application






ComputerHope.com



GPU

- GPU vs CPU



- GPU Cards

– GTX3090: **10496** cores, 24GB DDR6-384bit interface, Rs 1.3L



GPU Philosophy

- **Small independent function/code executed huge number of time**
- **Number of cores in thousands, tiny cores**
- **Cores are organized in cluster**
 - Kepler SMX: 14 SM, 192 SP Cuda cores/SM, 64 DP units, 32 SFU, 32 (LD/ST) U
 - TU102-RTX 2080Ti: 72 SM, 4608 Cuda Cores, 576 Tensor core, 72 RayTrace cores
- **Explicit memory hierarchy, programmer controlled**
- **Also implicit memory hierarchy : cache**

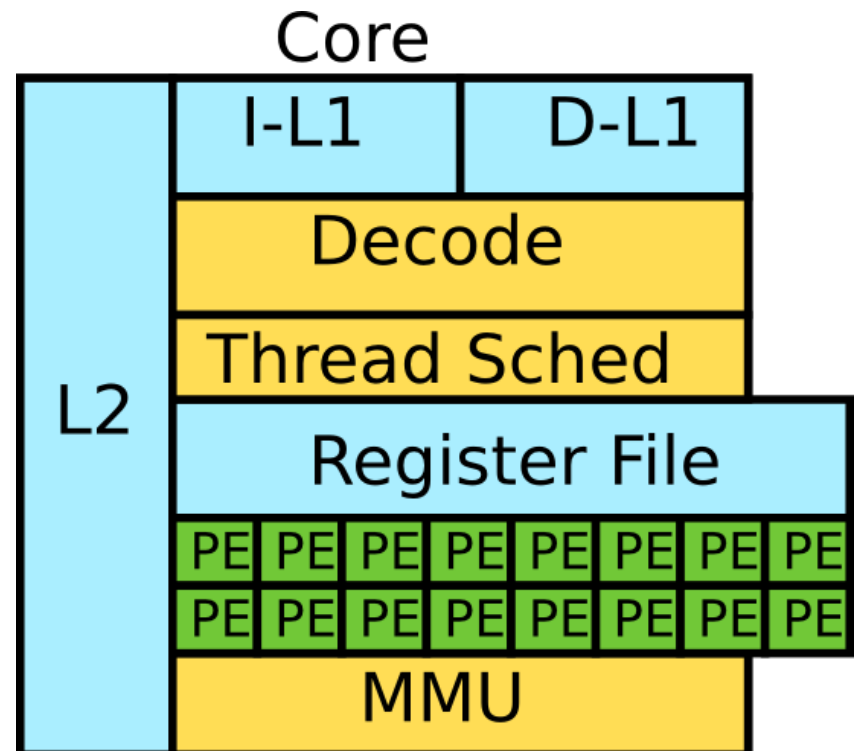
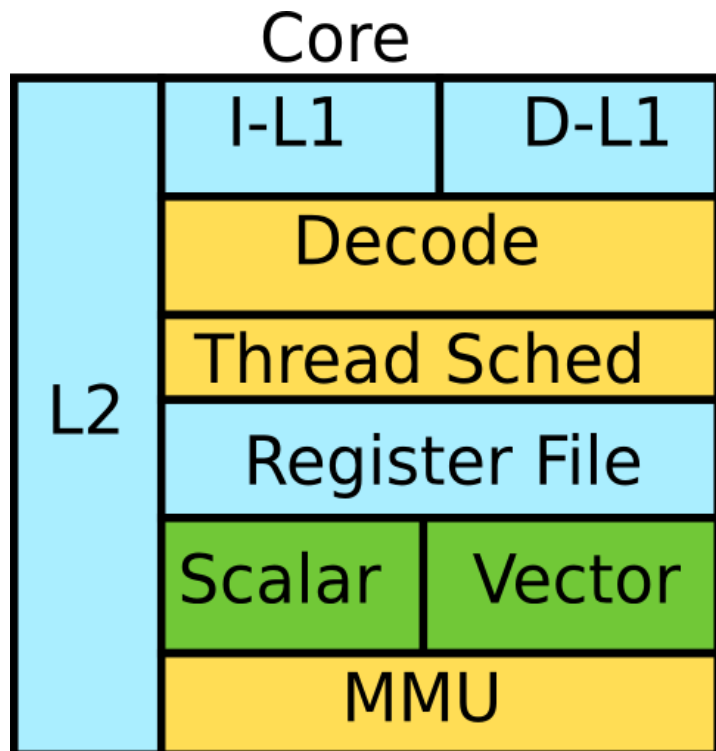
Single Instruction Multiple Data

GPU uses wide SIMD: 8/16/24/... processing elements (PEs)

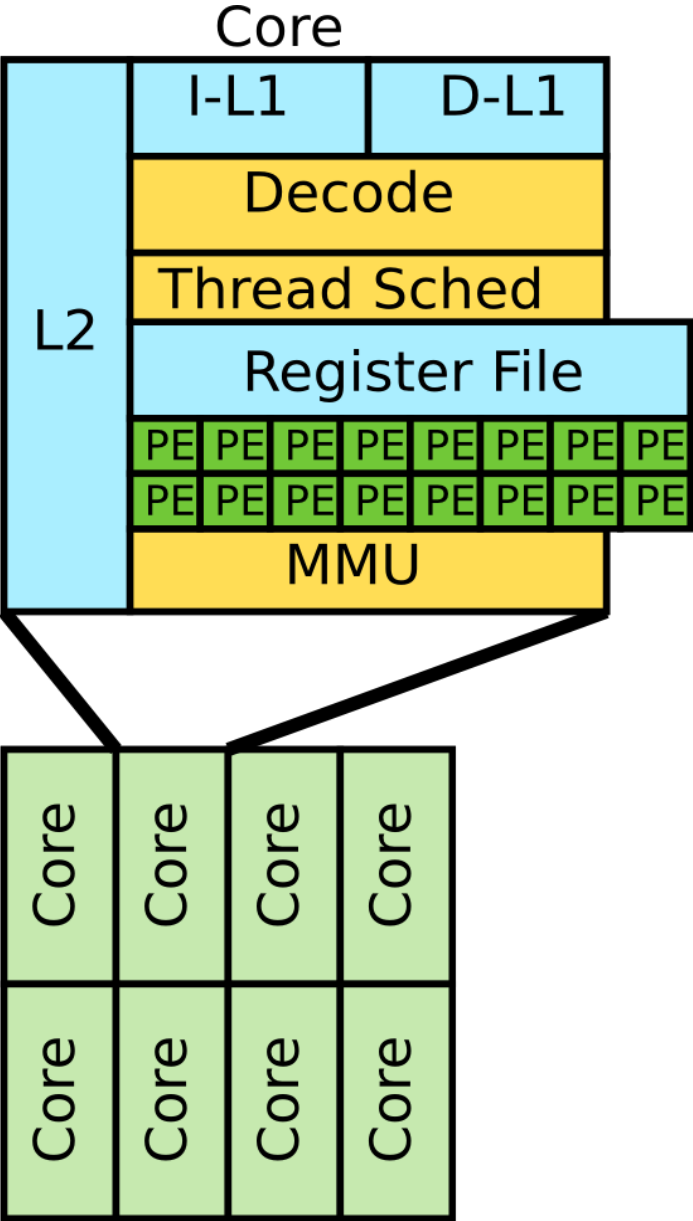
CPU uses short SIMD: usually has vector width of 4/8.

SSE has 4 data lanes

GPU has 8/16/24/... data lanes




Simple Core



The Stream Multiprocessor (SM) is a light weight core compared to IA core.



Stream Multiprocessor (SM)



Light weight PE:
Fused Multiply Add
(FMA)

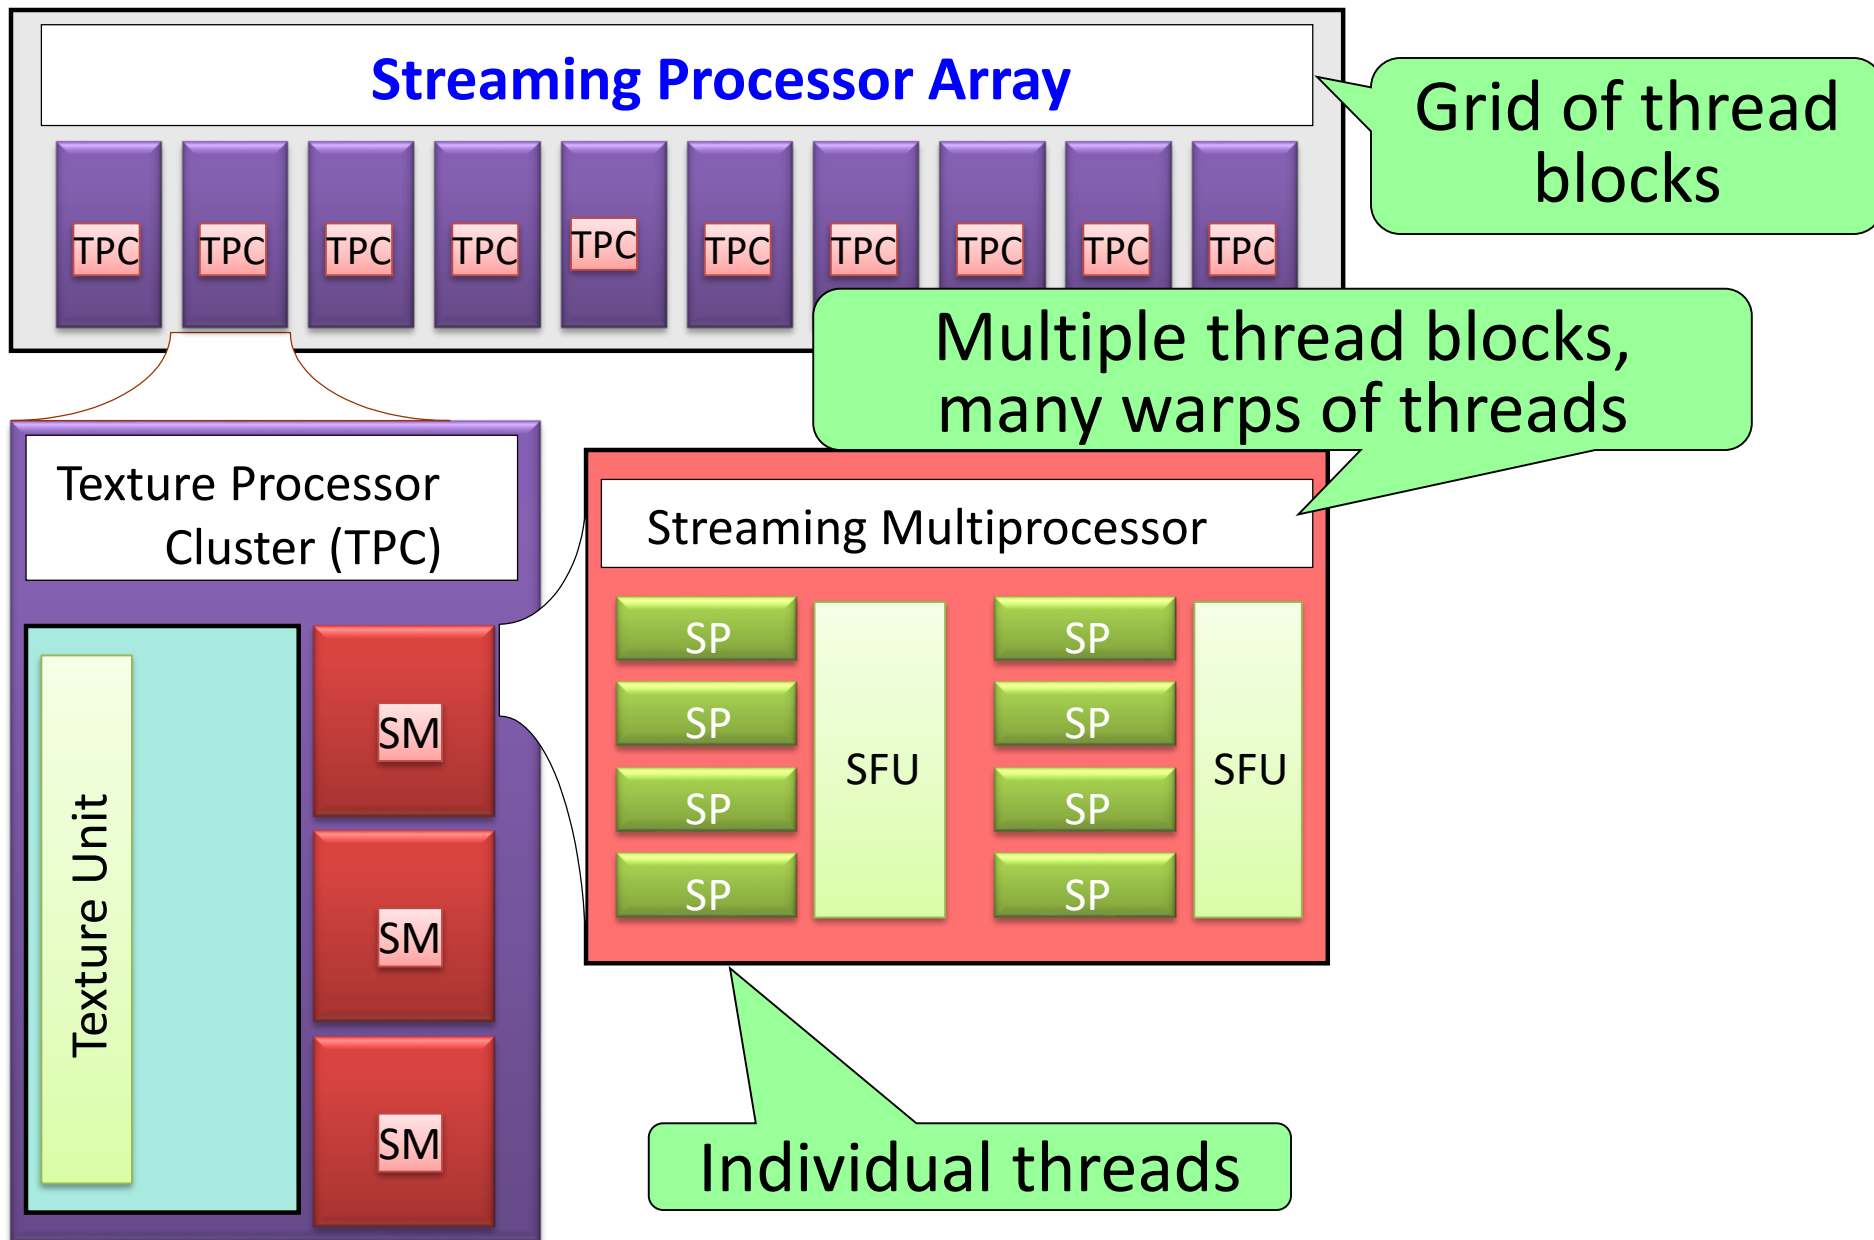
SFU:
Special Function
Unit

CPU vs GPU Layout



Source: Nvidia CUDA Programming Guide

NVIDIA architecture



GTX 690-Architecture



GeForce GTX 980Ti Technical Specs.


- 2816 cuda cores
- 22 SMX, 128 SP/SMX
- 6GB DDR5 RAM
- Each SMX have
 - Texture Cache 56KB
 - Constant memory (Scratch PAD) 65 KB
 - 49 KB L1/shared memory
 - Uniform Cache
 - Separate Shares Instruction Cache for SM

GeForce RTX 2080Ti Technical Specs.

- 4606 cuda cores
- 72 SMX, 64 SP/SMX
- **576 Tensor Core**
- **72 RayTrace Core**
- 11GB DDR6 RAM
- Each SMX have
 - Texture Cache 56KB
 - Constant memory (Scratch PAD) 65 KB
 - 49 KB L1/shared memory
 - Uniform Cache
 - Separate Shares Instruction Cache for SM

CPU-GPU Performance Gap

Theoretical GB/s



Source: CUDA Prog. Guide 4.0

Graphical Processing Units

- Given the hardware invested to do graphics well, how can be supplement it to improve performance of a wider range of applications?
- Basic idea:
 - Heterogeneous execution model
 - CPU is the *host*, GPU is the *device*
 - Develop a C-like programming language for GPU
 - Unify all forms of GPU parallelism as *CUDA thread*
 - Prog. model is “**Single Instruction Multiple Thread**”

Threads and Blocks

- A thread is associated with each data element
- Threads are organized into blocks
- Blocks are organized into a grid
- GPU hardware handles thread management, not applications or OS

Mapping SW Organized to GPU

- Software : Cuda Program or Cuda Threads
 - Thread organized into Grid, Block and Threads
 - Grid contain Blocks, Block contains many threads
- Hardware: GPU
 - Cores are organized in to clusters (SM)
 - GTX 690 per device: 8 SM, 192 core/SM
 - GTX980Ti: 22 SM, 128core/SM
- Mapping
 - Block get mapped to SM
 - Thread get map to core (SP)

Example thread scheduling

- Suppose we want to create parallel 2000 threads for an application
- We organize the thread into 10 blocks, each contains 200 threads
- When we run on top of GTX 690 GPU with 8 SM with 192 SP/SM
- Scheduler:
 - Map 10 blocks to 8 SM, takes $\text{ceil}(10/8) = 2$ times
 - Map 200 threads to 192 SP of SM, it also takes $\text{ceil}(200/192) = 2$ times
- **Thumb rules:**
 - Num block should be multiple of SM
 - Num thread/Block should be multiple of SP/SM

Example: DAXPY in C

```
//Invoke DAXPY
DAXPY (n, 2.0, x, y);
//function in C
void DAXPY (int n, double a,
            double *x, double *y) {
    for (int i=0; i<n; i++)
        y[i]=a*x[i]+y[i];
}
```

DAXPY in C and Cuda

```
//Invoke DAXPY with  
//256 threads per threadblock
```

```
__host__ int nb=(n+255)/256;  
DAXPY<<<nb,256>>>(n,2.0,x,y);
```

```
//function in Cuda
```

```
__device__ void DAXPY(int n, double a,  
                      double *x, double *y) {  
    int i=blockIdx.x*blockDim.x  
        +threadIdx.x;  
    if (i<n)  
        y[i]=a*x[i]+y[i];  
}
```

Complete DAXPY in Cuda

```
__device__ void DAXPY(int n, double a, double *x, double *y){
    int i;
    i=blockIdx.x*blockDim.x+threadIdx.x;
    if (i<n) y[i]=a*x[i]+y[i];
}


main(){
    #define N 1024
    int x[N], y[N], size=sizeof(int)*N;    Initialize(x,y);
    __device__ int *x_d, *y_d; //Declaration and Memory Creation
    cudaMalloc( (void **)&x_d, size); // On device
    cudaMalloc( (void **)&y_d, size);
    cudaMemcpy( x_d, x, size, cudaMemcpyHostToDevice );
    cudaMemcpy( y_d, y, size, cudaMemcpyHostToDevice );
    __host__ int nb=ceil(N/256.0); //Invoke DAXPY with 256 thrds/TB
    DAXPY<<<nb,256>>>(N,2.0,xd,yd);
    cudaMemcpy( y, y_d, size, cudaMemcpyDeviceToHost );
}
```

Two Levels of Thread Hierarchy

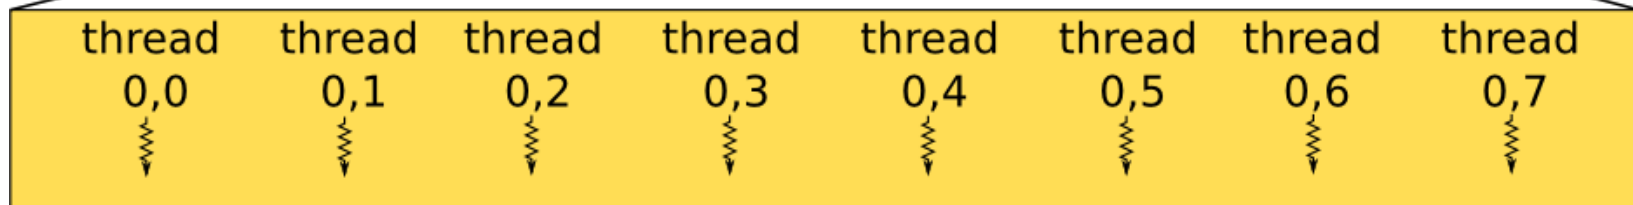
```
kernelF<<<(4,1),(8,1)>>>(A);
```

```
__device__ kernelF(A){  
    i = blockIdx.x;  
    j = threadIdx.x;  
    A[i][j]++;  
}
```

Grid kernelF contains 4 x 1 thread blocks



Thread Block



Each thread block contains 8 x 1 threads

Thread ⚡

Multi-dimension Thread and Block ID


Both grid and thread block can have two dimensional index.

```
kernelF<<<(2,2),(4,2)>>>(A);
```

```
__device__ kernelF(A){  
    i = blockDim.x * blockIdx.y  
      + blockIdx.x;  
    j = threadIdx.x * threadIdx.y  
      + threadIdx.x;  
    A[i][j]++;  
}
```


Grid

kernelF contains 2 x 2 thread blocks



Thread ⚡

Thread Block




Each thread block contains 4 x 2 threads


Scheduling Thread Blocks on SM

Grid

kernelF contains 2 x 2 thread blocks



Thread Block









Each thread block contains 4 x 2 threads

Example:

Scheduling 4 thread blocks on 3 SMs.

Thread ↕




Executing Thread Block on SM

```
kernelF<<<(2,2),(4,2)>>>(A);
```

```
__device__ kernelF(A){  
    i = blockDim.x * blockIdx.y  
        + blockIdx.x;  
    j = threadIdx.x * threadIdx.y  
        + threadIdx.x;  
    A[i][j]++;  
}
```


Thread Block




Thread 

Each thread block contains 4 x 2 threads

Executed on machine with width of 4:



Executed on machine with width of 8:




Examples: Application Mapping to Accelerator

- Machine Configuration
 - Pentium Quad core with 4GB RAM and installed GPU card with 3000 core and 4GB RAM on GPU
 - Need to consider Host to GPU memory data transfer and vice-verse.
- Applications
 - Vector Addition, Vector Sum
 - Matrix Multiplication
 - N-Body Simulation
 - Image Adaptive Histogram Equalization

Example 1: Vector Addition


```
int A[1000], B[1000], C[1000];  
for(i=0;i<1000;i++)  
  //Parallel, Independent Work  
  A[i]=B[i]+C[i];
```



- Executing on Host
 - $O(n)$ time, zero data transferred

Example 1: Vector Addition


```
int A[1000], B[1000], C[1000];  
for(i=0;i<1000;i++)  
//Parallel, Independent Work  
    A[i]=B[i]+C[i];
```



- Executing on GPU
 - 2×1000 data need to send from Host to GPU
 - Execution in $O(1)$ time parallel using 1000 cores
 - 1000 data need to be return from GPU to Host
- Take: 3000 unit communication overheads
funny, This **program is not a good candidate to run on GPU**

Example 1: Vector Addition


```
int A[1000], B[1000], C[1000];  
for(i=0;i<1000;i++)  
  //Parallel, Independent Work  
  A[i]=B[i]+C[i];
```



- Executing on the Host
 - using OpenMP/Pthread
 - Each thread can run from 250 locations
 - No need to transfer data, shared among all...
 - **Good news : (N/4 time)**

Example 2: Vector Sum


```
int A[1000];  
for(i=0;i<1000;i++)  
  //Parallel, Independent Work  
  S= S+A[i];
```



- Executing on Host
 - $O(n)$ time, zero data transferred

Example 2: Vector Sum


```
int A[1000];  
for(i=0;i<1000;i++)  
  //Parallel, Independent Work  
  S= S+A[i];
```



- Executing on GPU
 - 1000 data need to send from Host to GPU
 - Execution in $Lg(1000)$ time parallel using 1000 cores
(GPU core is have very weak support for shared variable locking)
 - 1 data need to be return from GPU to Host
- Take: 1000 unit communication overheads funny,
This **program is not a good candidate to run on GPU**

Example 2: Vector Sum


```
int A[1000];  
for(i=0;i<1000;i++)  
  //Parallel, Independent Work  
  S= S+A[i];
```



- Executing on the Host
 - Using OpenMP/Pthread
 - Each thread can run the sum from 250 locations
 - No need to transfer data, shared among all...
 - **Good news : time $(N/4 + 4)$**

Example 3: Matrix Multiplication


```
int A[100][100], B[100][100],  
C[100][100];  
for(i=0;i<100;i++)  
  for(j=0;j<100;j++){  
    Cij=0 ;  
    for(k=0;k<100;k++)  
      Cij=Cij+A[i][k]*B[k][j];  
  }
```



- Executing on Host
 - **$O(n^3)$ time**, zero data transferred
 - $O(n^{2.8})$ Strassen method

Example 3: Matrix Multiplication

```
int A[100][100], B[100][100],  
C[100][100];  
for(i=0;i<100;i++)  
  for(j=0;j<100;j++){  
    Cij=0 ;  
    for(k=0;k<100;k++)  
      Cij=Cij+A[i][k]*B[k][j];  
  }
```




Can be done in
one core

- Executing on GPU
 - $2 * n^2$ data send from Host to GPU + n^2 GPU to Host
 - $(n^2/3000) * n$ time parallel using 3000 cores
 - Significant reduction in running time
 - **This program is a good candidate to run on GPU**

Example 3: Matrix Multiplication

```
int A[100][100], B[100][100],  
C[100][100];  
for(i=0;i<100;i++)  
  for(j=0;j<100;j++){  
    Cij=0 ;  
    for(k=0;k<100;k++)  
      Cij=Cij+A[i][k]*B[k][j];  
  }
```



- Executing on the Host
 - Using OpenMP/Pthread
 - Each thread can run for $N^2/4$ Cij computation
 - No need to transfer data, shared among all...
 - **OK, not much reduction in time**

Example 4: N Body Simulation

```
typedef struct pos {float x,y,z;} pos;
typedef Force {float fx,fy,fz} Force;
PosP[N]; float Mass[N]; Force F[N];
while(1){
    for(i=0;i<N;i++){ //Calculate force for all body from others
        F[i]=0;
        for(j=0;j<N;j++) F[i]=F[i]+force([i][j]);
    }
    for(i=0;i<N;i++){ Ai=Fi/Mi; Pi=f(Ai,Pi);} //update Acc and Pos
}
}
```

- Executing on Host
 - $O(n^2)$ time for *each iteration of while loop...*,
zero data transferred , **for $N > 100$, Very slow**

Example 4: N Body Simulation

```
while(1){ //for each i, it can be in parallel
    for(i=0;i<N;i++){ //Calculate force for all body from others
        F[i]=0;
        for(j=0;j<N;j++)
            F[i]=F[i]+G*M[i]*M[j]/r[i][j]*r[i][j];
    } //for each i, it can be in parallel
    for(i=0;i<N;i++){
        A[i]=F[i]/M[i]; P[i]=f(A[i],P[i]);} //update A &P
    }
}
```

- Executing on GPU (**Embarrassingly parallel**)
 - 1*n data send from Host to GPU (old positions), 1*n from GPU to Host, onetime transfer of Mass
 - **Execution in $O(N)$ time parallel using N cores (if $N < 3000$)**
 - Total time: $1*n + O(N) + 1*n = O(n)$