

# **CS528**

# **Energy Efficient Task Scheduling**

A Sahu  
Dept of CSE, IIT Guwahati

# Outline

- Power and Energy Aware
- **Task with Hard Deadlines**
- Energy Efficiency
- Energy Efficient Scheduling
- **Real Time Task System**

# Top 500 HPC System

<https://www.top500.org/lists/top500/2021/11/>

## 1. Fugaku remains the No. 1 system. Japan

- 7,630,848 cores : **442 Pflop/s** on HPL Benchmarks.
- This puts it 3x ahead of the No. 2 system in the list.

## 2. Summit, an IBM-built system at the Oak Ridge National Laboratory (ORNL), USA,

- **148.8 Pflop/s** on the HPL benchmark
- 4,356 nodes, each housing
  - two Power9 CPUs with 22 cores each
  - and six NVIDIA Tesla V100 GPUs, each with 80 streaming multiprocessors (S.M.).
- Nodes are linked together with a Mellanox dual-rail EDR InfiniBand network.

## 3. Sierra, at Lawrence Livermore National Lab, USA,

- Architecture is very similar to Summit.
- 4,320 nodes with two Power9 CPUs and four NVIDIA Tesla V100 GPUs. Sierra achieved **94.6 Pflop/s**.


# Green500: evolution

- 2008: best result = **536 MFlops/Watt**
- 2009: best result = **723 MFlops/Watt**
  - Cell cluster, ranking 110 in top500
- 2010: best result = **1684 MFlops/Watt**
  - IBM BlueGene/Q prototype 1, ranking 101 in top500,
  - Peakperf: 65 TFlops;
- 2011: best result = **2097 MFlops/Watt**
  - IBM BlueGene/Q prototype 2
  - power consumption: 41 kW / Peak 85 TFlop/s

# Green500: evolution

<https://www.top500.org/>

Top 10 Green500 Supercomputers



# Problems of Energy Efficiency

- Laptop Problem
  - Given the energy budget, maximize number of Job
  - Given the Budget money maximize your satisfaction
    - Go to Restaurant with Rs 100. Choose Items to fill you stomach with your budget.
  - Given Rs 20 for going from IITG to Airport
    - Go to Jhalukbari using IIT G bus freely, Take another public bus pay Rs 20 to reach Airport.
  - Given Rs 10 : not possible, you need to walk...:)
  - Given Rs 600 how to go : Hire Taxi
  - Given Rs 20000 how to go : Hire BMW/Mercedes along with many other cars for security personals

# Problems of Energy Efficiency

- Server Problem
  - Budget is not constraints, minimize budget but do all the work (get all the items)
  - I want to Take all item of Thela/Bora..How much I need to pay? ---Bargaining

# Server Problem Example : $P_\infty | p_j, d_j | \Sigma E_j$

- We have infinite processors
- Processor can be run at speed  $f \in [0:1]$ ,  $PC = \alpha f^3$
- $N$  Tasks with deadlines, Task arrived at time 0, preemption not allowed,  $p_j$  at  $f=1$
- Execution time task  $t_j$  at freq  $f = e_j(t_j, f) = p_j / f$ ;
- Energy consumption task  $t_j$  at freq  $f$   
 $= E * \text{time} = PC(f) * e_j(t_j, f) = \alpha f^3 p_j / f = \alpha f^2 p_j$
- **We want to execute all the tasks, and minimize the sum of EC of all the tasks**



# Server Problem Example : $P_\infty | p_j, d_j | \Sigma E_j$

- We want to execute all the tasks, and minimize the sum of EC of all the tasks
- Solution
  - Select one processor for each of the tasks and total of  $N$  processors
  - Run the task at lowest feasible speed to meet the deadline  $f_j = p_j / d_j$
- This gives (optimal) minimum  $\Sigma E_j$ 
  - Total EC =  $\Sigma E_j = \Sigma \alpha f_j^2 p_j$
  - As  $(a+b)^2 > a^2 + b^2$  : running two task on one processor with higher speed consume higher energy

# Laptop Problem Example : $P_\infty, E_b \mid p_j, d_j \mid \Sigma U_j$

- We have infinite processors
- Processor can be run at speed  $f=[0:1]$ ,  $PC=\alpha f^3$
- N Tasks with deadlines, Task arrived at time 0, preemption not allowed,  $p_j$  at  $f=1$
- Execution time task  $t_j$  at freq  $f = e_j(t_j, f) = p_j / f$ ;
- Energy consumption task  $t_j$  at freq  $f$   
 $= E * \text{time} = PC(f) * e_j(t_j, f) = \alpha f^3 p_j / f = \alpha f^2 p_j$
- **We want to execute maximum number of the tasks before deadline given the energy budget**

# Laptop Problem Example : $P_{\infty}, E_b \mid p_j, d_j \mid \sum U_j$

- We want to execute maximum number of the tasks before deadline given the energy budget
- Solution:
  - Sort the tasks based on bare minimum energy requirement  $E_j = \alpha f_j^2 p_j$
  - Select the maximum number of task from this set
- Given N item with weight  $w_1, w_2, \dots, w_N$  : the weight is critical/min energy required of the task
- Select Maximum number of item given the Budget of Knapsack. **0-1 Knapsack Problem**
- NPC and Pseudo polynomial time algorithm exist using Dynamic Programming.

# Task with Deadline Vs Real Time Task System


- Task with Deadline:  $P | p_j | \Sigma U_j$ 
  - Every task have deadline
- Task with Soft Deadline
  - Deadline is not hard, but with QoS or Penalty
    - Airline provide *free sandwiches to flyer* when flight get delayed 1 hours/2 hours.
    - More than 3 hours of delay flyer are eligible for free cancellation
- Real time task system: every tasks occurs periodically
  - MP4: (a) video 30 F/S, (b) 16 bits, 2 Channel 44Khz
  - MP4: 1 video task, 2 audio tasks, : repeating one
- Soft Real time task system : Deadline can be soft

# Real Time Scheduling

- MPEG, Audio
  - 30 frame/Sec, 50 f/s, 60f/s
- Can you run 4K MKV file on Mobile ?
- Many Periodic Tasks in RT Systems
- **Nice Value** in **Linux**
  - 0-100 for real time task, 101-140 non real time task
  - Size of processor quantum (share) based on nice value

# Periodic Task: Real Time Scheduler

- Task with periods :  $T_i(c_i, p_i)$  here  $c_i$  is compute,  $p_i$ =period
- Each task have to finish before deadline with in the period



# Periodic Tasks

- Necessary schedulability test
  - Sum of utilization factors  $\mu_i$  must be less than or equal to  $n$ , where  $n$  is the number of processors
  - $\mu = \sum (c_i / p_i) \leq n$
  - $\mu_i$  = Percentage of time the task  $T_i$  requires the service of a CPU

# Periodic Task: Real Time Scheduler

## Assumptions & Definitions


- Tasks are periodic
- No aperiodic or sporadic tasks
- Job (instance) deadline = end of period
- Tasks are preemptable

- Laxity of a Task

$$T_i = d_i - (t + c_i')$$

where  $d_i$ : deadline;

$t$  : current time;       $c_i'$  : remaining computation time.






# Rate Monotonic Scheduling

- **Static Scheduling**

- Example: Suppose two RT tasks  $T_1(2,5)$  and  $T_2(3,7)$  :  $T_i(c_i, p_i)$  here  $c_i$  is compute,  $p_i$ =period
- Task with the smallest period is assigned the highest priority. At any time, the highest priority task is executed.



# Rate Monotonic (RM) Scheduling

- **Schedulability check (off-line)**

- A set of  $n$  tasks is schedulable on a uniprocessor by the RMS algorithm if the processor utilization (utilization test):

$$\sum_{i=1}^n \frac{c_i}{p_i} \leq n(2^{1/n} - 1)$$

The term  $n(2^{1/n} - 1)$  approaches  $\ln 2$ , ( $\approx 0.69$  as  $n \rightarrow \infty$ ).

# Earliest Deadline First (EDF)

- **Dynamic Scheduling**
- Task with the smallest deadline/laxity is assigned the highest priority. EDF or **Least Laxity First (LLF)**
  - At any time, the highest priority task is executed.
- **Schedulability check (off-line)**
  - A set of  $n$  tasks is schedulable on a uniprocessor by the EDF algorithm if the processor utilization.


$$\sum_{i=1}^n \frac{C_i}{P_i} \leq 1$$

- This condition is both necessary and sufficient.


# RM & EDF -- Example

Process	Period, $T$	WCET, $C$
$T_1$	5	2
$T_2$	7	4

EDF schedule



RMS schedule



# RT task: energy minimization

- Given a system of  $n$  periodic tasks  $T=\{\tau_1, \tau_2, \dots, \tau_n\}$  and one *Dynamic Volt-Freq Scaling Processor*
- With  $F=\{0, f_1, f_2, f_3, \dots, f_{max}\}$  finite number of freq
- And  $f_i < f_{i+1}$ .
  - Assume the task system satisfy  $\sum(wc_i/p_i) < 1$   
**at  $f_{max}$** ,  $wc_i$ =worst case compute time of  $i^{th}$  task
  - It ensure all the period task are schedulable without missing deadline if we run processor at  $f_{max}$


# RT task: energy minimization

- Design an efficient and elegant way to reduce the power/energy consumption ( $E=f^3t$ )
  - Hidden assumption: With out missing deadline of any task
- Number of processor is 1
- You may assume: all periodic tasks arrive at time 0
- Deadline of task is period of task

# Frequency Scaling EDF: Motivation

Pre-run schedule with holes

$WC_i$  = worst case computation time @  $F_{\max}$       Next arrival of T1



Holes in the pre-run schedule imply:

**EDF Test:**

$$\sum(wc_i/p_i) < 1 \quad \text{at frequency} = F_{\max}$$

In other words, whenever


**$\sum(wc_i/p_i) < 1$  there are holes in the EDF schedule**

# Frequency Scaling EDF: exploiting holes

Pre-run schedule with holes

$WC_i$  = **worst case computation time** @  $F_{\max}$

Next arrival  
of T1



Processor typically idles during holes. Instead, the holes can be exploited to slowdown the processor to save energy

How to do it ?

You need design an efficient and elegant way to reduce the Energy Consumption ?