

CS528

Rolling Horizon and Workload Prediction

A Sahu

Dept of CSE, IIT Guwahati

Outline

- Power Aware Scheduling in Cloud
- **Rolling Horizon**
- Work load Prediction

Energy Minimization Problem

- Given N tasks with (e_i, d_i) , $a_i=0$
- Given M machines with Power model of Machines

$$P(t) = P_{min} + \alpha \cdot u(t)^3$$

Execute all the tasks without missing deadlines and goal is to minimize energy.


Energy Model of Server/Host

- Processor is the major contributor
- Considered both static and dynamic EC

$$E = \int_{t_1}^{t_2} P(t) dt$$

$$P(t) = P_{min} + \alpha \cdot u(t)^3$$

Energy Model of Server/Host



- When a task runs with utilization u ,

$$E = (P_{min} + \alpha.u^3).e/u = (P_{min}/u + \alpha.u^2).e$$

Critical Utilization

Minimum utilization

- Utilization required by a task to finish at deadline

$$u_{i_min} = e_i / d_i$$

Least feasible VM

- Closest VM type which can support u_{i_min}

Critical utilization

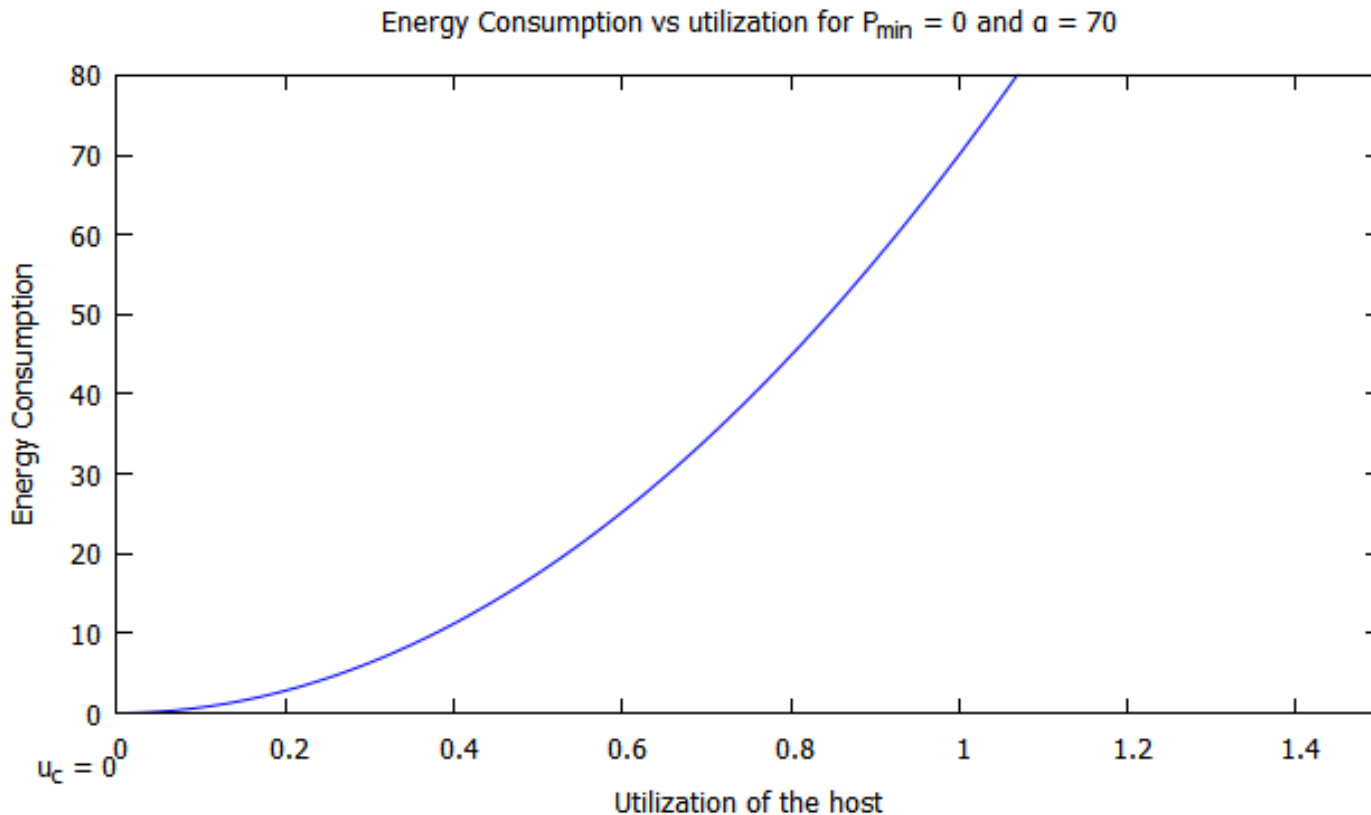
- Utilization of a host when energy consumption is minimum

$$u_c = \sqrt[3]{P_{min} / 2\alpha}$$

Classify of host based on U_c

- Case 1 : Negligible static power (i.e. $P_{min} = 0$)
 $\Rightarrow u_c = 0$
- Case 2 : Very high static power
 $\Rightarrow u_c > 1$
- Case 3 : General case
 $\Rightarrow 0 < u_c \leq 1$

Scheduling for Case I : $P_{\min}=0$

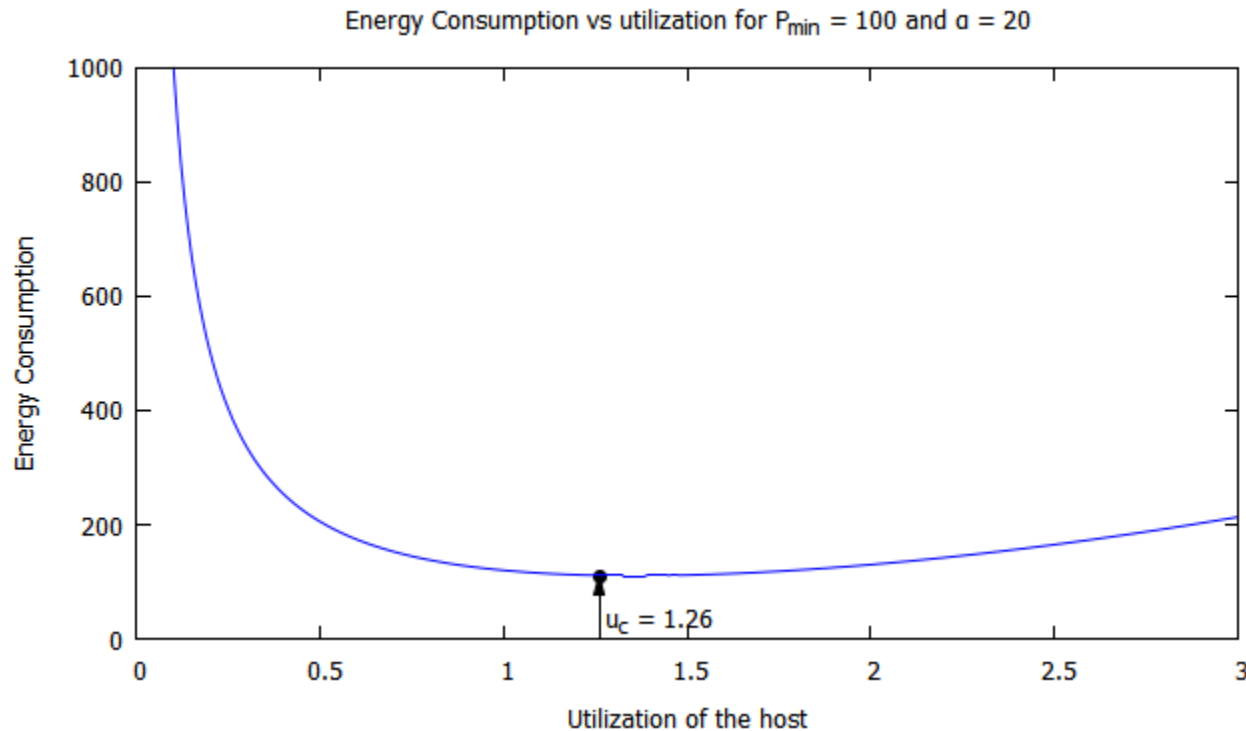


- For every task, choose the least feasible VM type
- Allocate the VM to a new host

$$E = \alpha \cdot u_t^2 \cdot e$$

$$u_1^2 + u_2^2 < (u_1 + u_2)^2$$


Scheduling for Case II : $U_c > 1$



- Maximum utilization of a running host is 1
- Closest feasible utilization value is 1
- Solved using general case with utilization 1

Energy Model of Server/Host

- Power consu. of a host, $P = P_{min} + \alpha u^3$



$$P_{min} = 100$$

$$\alpha = 70$$

- When a task runs with utilization u ,

$$E = (P_{min} + \alpha u^3) \cdot I / u = (P_{min} / u + \alpha u^2) \cdot I$$

Energy Model of Server/Host

Least feasible VM type

- A VM with utilization $\lceil u_{i_min} \rceil$
- $u_{i_min} = l_i / d_i$ (uti. to finish a task just at deadline)

Critical utilization

- utilization of a host when energy consumption is minimum

$$u_c = \sqrt[3]{P_{min} / 2\alpha}$$

**Target of the scheduler is to keep the
host utilization at u_c**

Reference

Zhu et. al, “Energy-Aware Rolling-Horizon Scheduling for Real-Time Tasks in Virtualized Cloud Data Centers”, IEEE HPCC 2013

Basic of Energy Aware RTS in Cloud

- Developing EA cloud data centers
 - not only can reduce power electricity cost
 - but also can improve system reliability.
- Real time task : task completion before deadline, QoS (weak term -reliable)
- Energy Saving using
 - Resource scaling up and scaling down strategies
- Rolling Horizon

Scheduling Model

- Given a virtualized data center that is characterized by an infinite set of physical computing hosts

$$H = \{h_1, h_2, \dots\}$$

- Hardware infrastructure H
 - for creating virtualized resources
 - to satisfy users' requirements.
- Active host set H_a with $H_a \subseteq H$.
- Host h_k is characterized by $h_k(c_k, r_k, n_k)$
 - Compute power in MIPS, ram size, net bandwidth

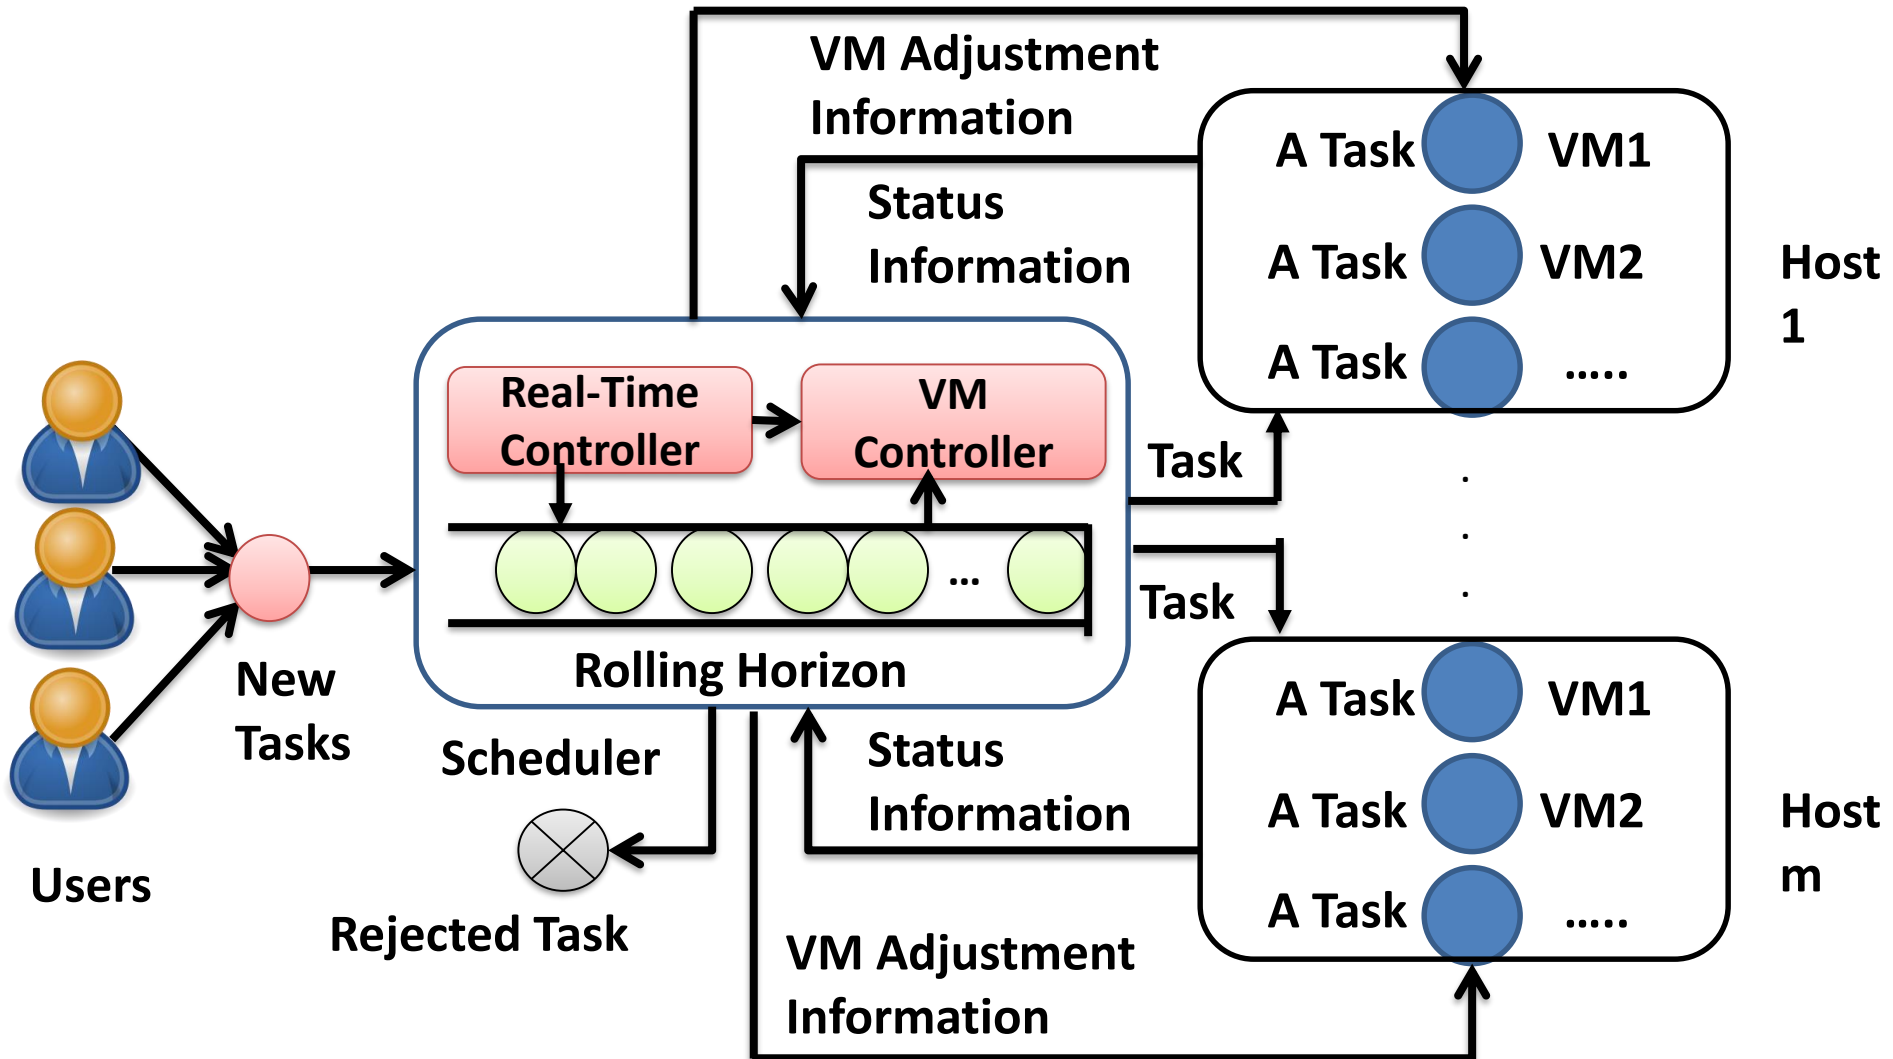
Scheduling Model

- Host h_k is have set V_k of virtual machine

$$V_k = \{ v_{1k}, v_{2k}, \dots, v_{|V_k|k} \}$$

- For a VM v_{jk} is characterized by $c(v_{jk})$, $r(v_{jk})$, $n(v_{jk})$
 - Fraction of Compute power in MIPS, ram size, net bandwidth allocated to v_{jk}
- Multiple VMs can be dynamically
 - Started and stopped on a single host
 - Based on the system workload.
- At the same time, some VMs are able
 - To migrate across hosts in order to consolidate resources
 - And further reduce energy consumption.

Scheduling Architecture



Working of RH Scheduler

- Scheduler consists of
 - Rolling-horizon, Real-time Controller, VM Controller.
- Scheduler work
 - Takes tasks from users and
 - Allocates them to different VMs.
- Rolling-horizon holds
 - Both new tasks and waiting tasks to be executed.
- A scheduling process is triggered
 - By new tasks, and all the tasks of RH to rescheduled

Working of RH Scheduler

- Scheduler consists of
 - Rolling-horizon, Real-time Controller, VM Controller.
- Scheduler work
 - Takes tasks from users and
 - Allocates them to different VMs.
- Rolling-horizon holds
 - Both new tasks and waiting tasks to be executed.
- A scheduling process is triggered
 - By new tasks, and all the tasks of RH to rescheduled

Scheduling ()@New Task Arrives

- **Step 1.** Scheduler checks System status information such as
 - running tasks' remaining execution time, active hosts, VMs' deployments,
 - Tasks in waiting pool including their deadlines,
 - Currently allocated VMs, start time, etc.
- **Step 2.** Sort the tasks in rolling-horizon
 - by their deadlines to facilitate scheduling operation.

Scheduling ()@New Task Arrives

- **Step 3.** Real-time controller determines
 - Whether a task in RH can be finished before its deadline.
- The VM controller adds VMs
 - to finish the task within timing constraint
 - if current VMs cannot finish it successfully.
 - If no schedule can be found to satisfy the task's timing requirement although enough VMs has been added by testing
- The task will be rejected. Or the task will be retained in the rolling-horizon.

Scheduling ()@New Task Arrives

- **Step 4.** Update the scheduling decision for the tasks in rolling-horizon,
 - Their execution order, start time,
 - Allocated VMs and new active hosts.
- **Step 5.** When a task in the rolling-horizon is ready to execute
 - dispatch the task to assigned VM.

Scheduling ()@New Task Arrives

- Additionally, when
 - tasks arrive slowly, tasks have loose deadlines or their count is less, making system workload light
- VM controller considers both
 - the status of active hosts and task information, and
- VM Controller then decides
 - Whether some VMs should be stopped or migrated to consolidate resources
 - So as to save energy

Task Model

- A set $T = \{t_1, t_2, \dots\}$ of independent tasks that arrive dynamically.
- A task t_i submitted by a user have
 $t_i = \{a_i, l_i, d_i, f_i\}$ Where a_i, l_i, d_i and f_i are
 - Arrival time, task length/size, deadline, and finish time of task t_i .
- Let rt_{jk} be the ready time of VM v_{jk} at host h_k .
- st_{ijk} be the start time of task t_i on VM v_{jk} .
- Execution time of task t_i on VM v_{jk} .

$$et_{ijk} = \frac{l_i}{c(v_{jk})}. \quad c(v): \text{compute capacity MIPS of VM}$$

Task Model

- Finish time of task t_i on v_{jk} , $ft_{ijk} = st_{ijk} + et_{ijk}$.
- Boolean x_{ijk} reflect mapping of tasks
 - to VMs at different hosts in a virtualized Cloud data center,

$x_{ijk} = 1$ if task t_i is allocated to VM v_{jk} at host h_k
 $= 0$, otherwise.

- Task's timing constraint can be guaranteed

$$x_{ijk} = \begin{cases} 0 & ft_{ijk} > d_i \\ 0 \text{ or } 1, & ft_{ijk} \leq d_i \end{cases}$$

Energy Consumption Model (ECM)

- EC by hosts in a data center
 - is mainly determined by CPU, memory, disk storage and network interfaces,
 - in which CPU consumes major part of energy
- CPU EC: static part (E_s) + dynamic part (E_d)
 - E_d is dominant > 80%, E_s follows similar trend to E_d
- EC of running task t_i on VM v_{jk} : $ec_{ijk} = ecr_{jk} \cdot et_{ijk}$
 - Term ecr_{jk} : EC rate of the VM v_{jk}

ECM of Tasks on VMs

- Total EC by executing all the tasks is

$$\begin{aligned} ec^{exec} &= \sum_{k=1}^{|H_a|} \sum_{j=1}^{|V_k|} \sum_{i=1}^T x_{ijk} \cdot ec_{ijk} \\ &= \sum_{k=1}^{|H_a|} \sum_{j=1}^{|V_k|} \sum_{i=1}^T x_{ijk} \cdot ecr_{jk} * et_{ijk} \end{aligned}$$

- H_a =active hosts, V_k =VM of host k , T =task set
- EC is incurred when VMs are sitting idle
 - All the VMs of a host is idle
 - Some of VM of a host Idle
- EC considering the execution time and idle time

$$ecei = ec^{exec} + ec^{allIdle} + ec^{partIdle}$$

ECM of host with idle VMs

- EC when all the VMs of a host is idle
 - Host can be set to a lower EC rate by DVFS
 - ECR of VM v_{jk} by ecr_{jk}^{idle}
 - Idle time when all the VMs in a host h_k are idle is it_k

$$ec^{allIdle} = \sum_{k=1}^{|Ha|} \sum_{j=1}^{|V_k|} ecr_{jk}^{idle} \cdot it_k$$

- EC when some of VM of a host Idle

$$ec^{partIdle} = \sum_{k=1}^{|Ha|} \sum_{j=1}^{|V_k|} ecr_{jk} \cdot t_j^{partIdle}$$

$$\text{with } t_j^{partIdle} = \max(f_i) - it_k - \sum_{i=1}^T x_{ijk} * et_{ijk}$$

Final ECM of host

- Although some VMs are placed on a host
 - maybe some resource is still unused
 - However, the resource also consume energy.
- Suppose there are s periods in each which the count of VMs in host h_k is different from another.
- Let t_p to denote the time in the period p ,
 $V_{k(p)}$ vm count in pth period on host k
- EC due to unused resources of hosts
 - $ec^{unused} = \sum_{k=1}^{|Ha|} \sum_{p=1}^s \left(ecr(h_k) - \sum_{j=1}^{V_{k(p)}} ecr_{jk} \right) \cdot tp$
- So total EC of Cloud system

$$ec = ecei + ec^{unused} = ec^{exec} + ec^{allidle} + ec^{partidle} + ec^{unused}$$

Scheduling Goals and Trade-offs

- Less running hosts ==>
 - less consumed energy
 - may greatly affect **guarantee ratio** of real-time **tasks**
- Energy Conservation and TGR are two conflicting objectives
- Good scheduling strategy makes a good trade-off by dynamically
 - Starting hosts, Closing hosts
 - Creating VMs, canceling VMs
 - Migrating VMs
 - **according to the system workload.**

Energy Aware Rolling Horizon Scheduling

- Traditional scheduling: once a task is scheduled
 - it is dispatched immediately to the local queue of a VM or a host
- In EARH: puts all the waiting tasks RH queue
 - Their schedules are allowed to be adjusted for the schedulability of the new task
 - possibly less energy consumption.
- Essential advantage of RH optimization
 - task migration required by rescheduling does not yield any overhead
 - as all the tasks are waiting in the rolling-horizon.

EARH Approach

- Attempt to append new task to the end of former allocated tasks on a VM.
- Start time st_{ijk} of task t_i on VM v_{jk}
 $st_{ijk} = \max\{rt_{jk}, a_i\}$, rt_{jk} is ready time of v_{jk}
- Ready time get updated once task t_p added to v_{jk}
 $rt_{jk} = st_{pjk} + et_{pjk}$.
- When a task cannot be successfully allocated in any current VM
 - `scaleUpResource()` is done to create a new VM
 - With the goal of finishing the task within its deadline.

Scale up resources: in three steps

- Step1: Create a new VM
 - in a current active host without any VM migration
- Step2: If Step 1 fails
 - Migrate some VMs among current active hosts
 - To yield enough resource on a host and
 - then create a new VM on it
- Step3: if Step 2 fails
 - start a host and then create a new VM on it.

Scale up resources: in three steps

- *Some terms*

- $st(h_k)$: start-up time of host h_k
- $ct(v_{jk})$: creation time of VM v_{jk}
- $mt(v_{jk})$: migration time of VM v_{jk} , $=r(v_{jk})/n(v_{jk})$,
RAM and Network

- Using different steps products different start times for a task,

$$st_{ijk} = \begin{cases} a_i + ct(v_{jk}), & \text{if setp1,} \\ a_i + ct(v_{jk}) + \sum_{p=1}^{|p|} mt(v_{pk}), & \text{if setp2,} \\ a_i + st(h_k) + ct(v_{jk}), & \text{if setp3.} \end{cases}$$

Scheduling EARH

for each task t_i in set Q **do**

$findTag \leftarrow \text{FALSE}; findVM \leftarrow \text{NULL};$

for each VM v_{jk} in the system **do**

Calculate the start time st_{ijk} and execution time et_{ijk}

If $st_{ijk} + et_{ijk} \leq d_i$ **then** $findTag \leftarrow \text{TRUE};$ Compute ec_{ijk}

if $findTag == \text{FALSE}$ **scaleUpResource();**

if $findTag == \text{TRUE}$ **then**

Select v_{sk} with min energy consumption to execute t_i

$findVM \leftarrow v_{sk}$

else Reject task t_i

Update scheduling decision of t_i and remove it from Q